

QUESTION 883

Question : A solutions architect is responsible for building an architecture that will support the operation of a third-party database server. The database software is memory heavy and is licensed on a CPU-based basis, with the cost increasing in direct proportion to the number of virtual CPU cores in the operating system. The solutions architect must choose an Amazon EC2 instance with adequate RAM to operate the database software, yet with a high number of vCPUs. The solutions architect must guarantee that the virtual CPUs are not underutilized and must keep expenditures to a minimum. Which solution satisfies these criteria ?

- A.Select and launch a smaller EC2 instance with an appropriate number of vCPUs.
- B.Configure the CPU cores and threads on the selected EC2 instance during instance launch.
- C.Create a new EC2 instance and ensure multithreading is enabled when configuring the instance details.
- D.Create a new Capacity Reservation and select the appropriate instance type. Launch the instance into this new Capacity Reservation.

Answer : B

Explanation : B-Specify a custom number of vCPUs (while enjoying the same memory, storage, and bandwidth of a full-sized instance) when launching new Amazon RDS for Oracle instances.

QUESTION 884

Question : A solutions architect is creating a two-tiered architecture with distinct private subnets for compute and database resources. AWS Lambda functions deployed in compute subnets need database connection. Which option would provide the MOST SECURE connectivity ?

- A.Configure the Lambda function to use Amazon RDS Proxy outside the VPC.
- B.Associate a security group with the Lambda function. Authorize this security group in the database's security group.
- C.Authorize the compute subnet's CIDR ranges in the database's security group.
- D.During the initialization phase, authorize all IP addresses in the database's security group temporarily. Remove the rule after the initialization is complete.

Answer : C

Explanation : Authorize the compute subnets CIDR ranges in the database's security

group Because, we need to, make sure that For the Lambda function's security group, that traffic is allowed to go in and out of the CIDR of the RDS instance's VPC.

QUESTION 885

Question : A business wishes to use a customized distributed program for the purpose of calculating numerous profit and loss situations.To do this, the business must establish a network connection between its Amazon EC2 instances.The connection must have a lowlatency and a high throughput.Which solution will satisfy these criteria ?

- A.Provision the application to use EC2 Dedicated Hosts of the same instance type.
- B.Configure a placement group for EC2 instances that have the same instance type.
- C.Use multiple AWS elastic network interfaces and link aggregation.
- D.Configure AWS PrivateLink for the EC2 instances.

Answer : B

Explanation : Cluster Placement group.While launching select instance type

QUESTION 886

Question : A solution architect is designing a new workload in which AWS lambda function will access an amazon DynamoDB table.what is most secure means of granting the lambda function access to the DynamoDB table ?

- A.Create a IAM role with necessary permissions to access the DynamoDB.Assign the role to lambda function.
- B.Create Dynamo username and password and give them to developer to use in the lambda function.
- C.Create an IAM user and create access and secret keys for the user.Give the user necessary permissions to access DynamoDB table
- D.Create a IAM role allowing

Answer : D

QUESTION 887

Question : Every 90 days, a security team must enforce the rotation of all IAM users'access keys.If an access key is discovered to be older, it must be disabled and deleted.A solutions architect must design a solution that will detect and remediate keys

that are more than 90 days old.Which method satisfies these criteria with the LEAST amount of operational effort ?

- A.Create an AWS Config rule to check for the key age.Configure the AWS Config rule to run an AWS Batch job to remove the key.
- B.Create an Amazon EventBridge (Amazon CloudWatch Events) rule to check for the key age.Configure the rule to run an AWS Batch job to remove the key.
- C.Create an AWS Config rule to check for the key age.Define an Amazon EventBridge (Amazon CloudWatch Events) rule to schedule an AWS Lambda function to remove the key.
- D.Create an Amazon EventBridge (Amazon Cloud Watch Events) rule to check for the key age.Define an EventBridge (Cloud Watch Events) rule to run an AWS Batch job to remove the key.

Answer : A.

Explanation :

https://docs.aws.amazon.com/service-authorization/latest/reference/list_awscopy.html

QUESTION 888

Question : A company collects 10 GB of telemetry data dairy from various machines.The company stores the data in an Amazon S3 bucket in a source data account.The company has hired several consuming agencies to use this data for analysis.Each agency needs read access to the data for its analysis.The company must share the data from the source data account by choosing a solution that maximizes security and operational efficiency.

Which solution will meet these requirements ?

- A.Configure S3 global tables to replicate data tor each agency
- B.Make the S3 bucket public for a limited time Inform only the agencies
- C.Configure cross-account access for the S3 bucket to the accounts that the agencies own.
- D.Set up an IAM user for each analyst In the source data account Grant each user access to the S3 bucket

Answer : C

QUESTION 889

Question : Prior to delivering a new workload, a solutions architect must examine and

update the organization's current IAM rules. The following policy was written by the solutions architect: What is the policy's net effect?

- A. Users will be allowed all actions except s3 : PutObject if multi-factor authentication (MFA) is enabled.
- B. Users will be allowed all actions except s3 : PutObject if multi-factor authentication (MFA) is not enabled.
- C. Users will be denied all actions except s3 : PutObject if multi-factor authentication (MFA) is enabled.
- D. Users will be denied all actions except s3 : PutObject if multi-factor authentication (MFA) is not enabled.

Answer : D

QUESTION 890

Question : A company is deploying an application that processes large quantities of data in batches as needed. The company plans to use Amazon EC2 instances for the workload. The network architecture must support a highly scalable solution and prevent groups of nodes from sharing the same underlying hardware. Which combination of network solutions will meet these requirements? (Select TWO.)

- A. Create Capacity Reservations for the EC2 instances to run in a placement group.
- B. Run the EC2 instances in a spread placement group.
- C. Run the EC2 instances in a cluster placement group.
- D. Place the EC2 instances in an EC2 Auto Scaling group.
- E. Run the EC2 instances in a partition placement group.

Answer : BE

Explanation : Spread-strictly places a small group of instances across distinct underlying hardware to reduce correlated failures. Partition-spreads your instances across logical partitions such that groups of instances in one partition do not share the underlying hardware with groups of instances in different partitions. This strategy is typically used by large distributed and replicated workloads, such as Hadoop, Cassandra, and Kafka.

QUESTION 891

Question : A company is designing a new web application that the company will deploy into a single AWS Region. The application requires a two-tier architecture that will include Amazon EC2 instances and an Amazon RDS DB instance. A solutions architect needs to design the application so that all components are highly available. Which solution will meet these requirements MOST cost-effectively ?

- A. Deploy EC2 instances in an additional Region. Create a DB instance with the Multi-AZ option activated.
- B. Deploy all EC2 instances in the same Region and the same Availability Zone. Create a DB instance with the Multi-AZ option activated.
- C. Deploy EC2 instances across at least two Availability Zones within the same Region. Create a DB instance in a single Availability Zone.
- D. Deploy EC2 instances across at least two Availability Zones within the same Region. Create a DB instance with the Multi-AZ option activated.

Answer : D

Explanation : App will be deployed in single AWS region. Hence A is wrong. B-instances in same AZ i.e. no HA. C-DB is not HA but as per the QUESTION all components must be highly available. Hence, incorrect. D-fulfills all the requirements.

QUESTION 892

Question : A company has migrated a two-tier application from its on-premises data center to the AWS Cloud. The data tier is a Multi-AZ deployment of Amazon RDS for Oracle with 12c of General Purpose SSD Amazon Elastic Block Store (Amazon EBS) storage. The application is designed to process and store documents in the database as binary largeobjects (blobs) with an average document size of 6 MB. The database size has grown over time, reducing the performance and increasing the cost of storage. The company must improve the database performance and needs a solution that is highly available and resilient. Which solution will meet these requirements MOST cost-effectively ?

- A. Reduce the RDS DB instance size. Increase the storage capacity to 24 TiB. Change the storage type to Magnetic.
- B. Increase the RDS DB instance size. Increase the storage capacity to 24 TiB. Change the storage type to Provisioned IOPS.
- C. Create an Amazon S3 bucket. Update the application to store documents in the S3

bucket.Store the object metadata in the existing database.

D.Create an Amazon DynamoDB table.Update the application to use DynamoDB.Use AWS Database Migration Service (AWS DMS) to migrate data from the Oracle database to DynamoDB.

Answer : C

Explanation : Because of Binary large object (BLOB) storage.DynamoDB can store binary items up to 400 KB, but DynamoDB is not generally suited to storing documents or images.A better architectural pattern for this implementation is to store pointers to Amazon S3 objects in a DynamoDB table

QUESTION 893

Question : A company runs an application on a large fleet of Amazon EC2 instances.The application reads and writes entries into an Amazon DynamoDB table.The size of the DynamoDB table continuously grows, but the application needs only data from the last 30

days.The company needs a solution that minimizes cost and development effort.Which solution meets these requirements ?

A.Use an AWS CloudFormation template to deploy the complete solution.Redeploy the CloudFormation stack every 30 days, and delete the original stack.

B.Use an EC2 instance that runs a monitoring application from AWS Marketplace.Configure the monitoring application to use Amazon DynamoDB Streams to store the timestamp when a new item is created in the table.Use a script that runs on the EC2 instance to delete items that have a timestamp that is older than 30 days.

C.Configure Amazon DynamoDB Streams to invoke an AWS Lambda function when a new item is created in the table.Configure the Lambda function to delete items in the table that are older than 30 days.

D.Extend the application to add an attribute that has a value of the current timestamp plus

30 days to each new item that is created in the table.Configure DynamoDB to use the attribute as the TTL attribute.

Answer : D

Explanation : Because the QUESTION is asking retain data for 30 days.Using TTL attribute is best option here.A Wrong, nothing to do with cfn.B Wrong, no need to use EC2 C Wrong, no need to use lambda.

QUESTION 894

Question : A company has many projects that run in multiple AWS Regions. The projects usually have a three-tier architecture with Amazon EC2 instances that run behind an Application Load Balancer. The instances run in an Auto Scaling group and share Amazon Elastic File System (Amazon EFS) storage and Amazon RDS databases. Some projects have resources in more than one Region. A solutions architect needs to identify each project's individual costs. Which solution will provide this information with the LEAST amount of operational effort ?

- A.Use Cost Explorer to perform one-time queries for each Region and create a report that filters by project.
- B.Use the AWS Billing and Cost Management details page to see the actual usage costs of the resources by project.
- C.Use AWS Systems Manager to group resources by project and monitor each project's resources and cost.
- D.Use AWS Billing and Cost Management to activate cost allocation tags and create reports that are based on the project tags.

Answer : D

Explanation : Using Cost Allocation Tags=the LEAST amount of operational effort

QUESTION 895

Question : A company wants to perform an online migration of active datasets from an on-premises NFS server to an Amazon S3 bucket that is named DOC-EXAMPLE-BUCKET. Data integrity verification is required during the transfer and at the end of the transfer. The data also must be encrypted. A solutions architect is using an AWS solution to migrate the data. Which solution meets these requirements ?

- A.AWS Storage Gateway file gateway
- B.S3 Transfer Acceleration
- C.AWS DataSync
- D.AWS Snowball Edge Storage Optimized

Answer : C

Explanation : because AWS DataSync handle validating data integrity

QUESTION 896

Question : A news company that has reporters all over the world is hosting its broadcast system on AWS. The reporters send live broadcasts to the broadcast system. The reporters use software on their phones to send live streams through the Real Time Messaging Protocol (RTMP) . A solutions architect must design a solution that gives the reporters the ability to send the highest quality streams. The solution must provide accelerated TCP connections back to the broadcast system. What should the solutions architect use to meet these requirements ?

- A.Amazon CloudFront
- B.AWS Global Accelerator
- C.AWS Client VPN
- D.Amazon EC2 instances and AWS Elastic IP addresses

Answer : B

Explanation : AWS Global Accelerator and Amazon CloudFront are separate services that use the AWS global network and its edge locations around the world.

CloudFront improves performance for both cacheable content (such as images and videos) and dynamic content (such as API acceleration and dynamic site delivery) . Global Accelerator improves performance for a wide range of applications over TCP or UDP by proxying packets at the edge to applications running in one or more AWS Regions. Global Accelerator is a good fit for non-HTTP use cases, such as gaming (UDP) , IoT (MQTT) or Voice over IP, as well as for HTTP use cases that specifically require static IP addresses or deterministic, fast regional failover.

QUESTION 897

Question : A company is designing a shared storage solution for a gaming application that is hosted in the AWS Cloud. The company needs the ability to use SMB clients to access data. The solution must be fully managed. Which AWS solution meets these requirements ?

- A.Create an AWS DataSync task that shares the data as a mountable file system. Mount the file system to the application server.
- B.Create an Amazon EC2 Windows instance. Install and configure a Windows file share role on the instance. Connect the application server to the file share.
- C.Create an Amazon FSx for Windows File Server file system. Attach the file system

to the origin server.Connect the application server to the file system.

D.Create an Amazon S3 bucket.Assign an IAM role to the application to grant access to the S3 bucket.Mount the S3 bucket to the application server.

Answer : C

Explanation : since it is an SMB file store (for Windows) +Gaming needs

HPC

QUESTION 898

Question : A solutions architect needs to design the architecture for an application that a vendor provides as a Docker container image.The container needs 50 GB of storage available for temporary files.The infrastructure must be serverless.Which solution meets these requirements with the LEAST operational overhead ?

A.Create an AWS Lambda function that uses the Docker container image with an Amazon S3 mounted volume that has more than 50 GB of space.

B.Create an AWS Lambda function that uses the Docker container image with an Amazon Elastic Block Store (Amazon EBS) volume that has more than 50 GB of space.

C.Create an Amazon Elastic Container Service (Amazon ECS) cluster that uses the AWS Fargate launch type.Create a task definition for the container image with an Amazon Elastic File System (Amazon EFS) volume.Create a service with that task definition.

D.Create an Amazon Elastic Container Service (Amazon ECS) cluster that uses the Amazon EC2 launch type with an Amazon Elastic Block Store (Amazon EBS) volume that has more than 50 GB of space.Create a task definition for the container image.Create a service with that task definition.

Answer : C

Explanation : AWS Fargate is a serverless, pay-as-you-go compute engine that lets you focus on building applications without managing servers.AWS Fargate is compatible with both Amazon Elastic Container Service (ECS) and Amazon Elastic Kubernetes Service (EKS)

QUESTION 899

Question : A company wants to share data that is collected from self-driving cars with the automobile community.The data will be made available from within an Amazon S3

bucket. The company wants to minimize its cost of making this data available to other AWS accounts. What should a solutions architect do to accomplish this goal ?

- A.Create an S3 VPC endpoint for the bucket.
- B.Configure the S3 bucket to be a Requester Pays bucket.
- C.Create an Amazon CloudFront distribution in front of the S3 bucket.
- D.Require that the files be accessible only with the use of the Bit Torrent protocol.

Answer : B

QUESTION 900

Question : A company is running several business applications in three separate VPCs within the eu-east-1 Region. The applications must be able to communicate between VPCs. The applications also must be able to consistently send hundreds of gigabytes of data each day to a latency-sensitive application that runs in a single on-premises data center. A solutions architect needs to design a network connectivity solution that maximizes cost-effectiveness. Which solution meets these requirements ?

- A.Configure three AWS Site-to-Site VPN connections from the data center to AWS. Establish connectivity by configuring one VPN connection for each VPC.
- B.Launch a third-party virtual network appliance in each VPC. Establish an IPsec VPN tunnel between the data center and each virtual appliance.
- C.Set up three AWS Direct Connect connections from the data center to a Direct Connect gateway in us-east-1. Establish connectivity by configuring each VPC to use one of the Direct Connect connections.
- D.Set up one AWS Direct Connect connection from the data center to AWS. Create a transit gateway, and attach each VPC to the transit gateway. Establish connectivity between the Direct Connect connection and the transit gateway.

Answer : D

QUESTION 901

Question : A company is hosting a three-tier ecommerce application in the AWS Cloud. The company hosts the website on Amazon S3 and integrates the website with an API that handles sales requests. The company hosts the API on three Amazon EC2 instances behind an Application Load Balancer (ALB) . The API consists of static and dynamic front-end content along with backend workers that process sales requests asynchronously. The company is expecting a significant and sudden increase in

the number of sales requests during events for the launch of new products.What should a solutions architect recommend to ensure that all the requests are processed successfully ?

- A.Add an Amazon CloudFront distribution for the dynamic content.Increase the number of EC2 instances to handle the increase in traffic.
- B.Add an Amazon CloudFront distribution for the static content.Place the EC2 instances in an Auto Scaling group to launch new instances based on network traffic.
- C.Add an Amazon CloudFront distribution for the dynamic content.Add an Amazon ElastiCache instance in front of the ALB to reduce traffic for the API to handle.
- D.Add an Amazon CloudFront distribution for the static content.Add an Amazon Simple Queue Service (Amazon SQS) queue to receive requests from the website for later processing by the EC2 instances.

Answer : D

QUESTION 902

Question : A company is designing an application where users upload small files into Amazon S3.After a user uploads a file, the file requires one-time simple processing to transform the data and save the data in JSON format for later analysis.Each file must be processed as quickly as possible after it is uploaded.Demand will vary.On some days, users will upload a high number of files.On other days, users will upload a few files or no files.Which solution meets these requirements with the LEAST operational overhead ?

- A.Configure Amazon EMR to read text files from Amazon S3.Run processing scripts to transform the data.Store the resulting JSON file in an Amazon Aurora DB cluster.
- B.Configure Amazon S3 to send an event notification to an Amazon Simple Queue Service (Amazon SQS) queue.Use Amazon EC2 instances to read from the queue and process the data.Store the resulting JSON file in Amazon DynamoDB.
- C.Configure Amazon S3 to send an event notification to an Amazon Simple Queue Service (Amazon SQS) queue.Use an AWS Lambda function to read from the queue and process the data.Store the resulting JSON file in Amazon DynamoDB.
- D.Configure Amazon EventBridge (Amazon CloudWatch Events) to send an event to

Amazon Kinesis Data Streams when a new file is uploaded. Use an AWS Lambda function to consume the event from the stream and process the data. Store the resulting JSON file in Amazon Aurora DB cluster.

Answer : C

QUESTION 903

Question : An image-hosting company stores its objects in Amazon S3 buckets. The company wants to avoid accidental exposure of the objects in the S3 buckets to the public. All S3 objects in the entire AWS account need to remain private. Which solution will meet these requirements ?

- A. Use Amazon GuardDuty to monitor S3 bucket policies. Create an automatic remediation action rule that uses an AWS Lambda function to remediate any change that makes the objects public.
- B. Use AWS Trusted Advisor to find publicly accessible S3 buckets. Configure email notifications in Trusted Advisor when a change is detected. Manually change the S3 bucket policy if it allows public access.
- C. Use AWS Resource Access Manager to find publicly accessible S3 buckets. Use Amazon Simple Notification Service (Amazon SNS) to invoke an AWS Lambda function when a change is detected. Deploy a Lambda function that programmatically remediates the change.
- D. Use the S3 Block Public Access feature on the account level. Use AWS Organizations to create a service control policy (SCP) that prevents IAM users from changing the setting. Apply the SCP to the account.

Answer : D

Explanation : While guard duty helps to monitor s3 for potential threats its a reactive action. We should always be proactive and not reactive in our solutions so D, block public access to avoid any possibility of the info becoming publicly accessible.

QUESTION 904

Question : A company has a website hosted on AWS. The website is behind an Application Load Balancer (ALB) that is configured to handle HTTP and HTTPS separately. The company wants to forward all requests to the website so that the requests will use HTTPS. What solution should a solutions architect do to meet this requirement ?

- A. Update the ALB's network ACL to accept only HTTPS traffic.

- B.Create a rule that replaces the HTTP in the URL with HTTPS.
- C.Create a listener rule on the ALB to redirect HTTP traffic to HTTPS.
- D.Replace the ALB with a Network Load Balancer configured to use Server Name Indication (SNI)

Answer : C

QUESTION 905

Question : A company used an AWS Direct Connect connection to copy 1 PB of data from a colocation facility to an Amazon S3 bucket in the us-east-1 Region. The company now wants to copy the data to another S3 bucket in the us-west-2 Region. Which solution will meet this requirement ?

- A.Use an AWS Snowball Edge Storage Optimized device to copy the data from the colocation facility to us-west-2.
- B.Use the S3 console to copy the data from the source S3 bucket to the target S3 bucket.
- C.Use S3 Transfer Acceleration and the S3 copy-object command to copy the data from the source S3 bucket to the target S3 bucket.
- D.Add an S3 Cross-Region Replication configuration to copy the data from the source S3 bucket to the target S3 bucket.

Answer : D

Explanation : Depending on your use case, you can perform the data transfer between buckets using one of the following options: Run parallel uploads using the AWS Command Line Interface (AWS CLI) ; Use an AWS SDK ; Use cross-Region replication or same-Region replication ; Use Amazon S3 batch operations ; Use S3DistCp with Amazon EMR ; Use AWS DataSync.

QUESTION 906

Question : A company recently announced the deployment of its retail website to a global audience. The website runs on multiple Amazon EC2 instances behind an Elastic Load Balancer. The instances run in an Auto Scaling group across multiple Availability Zones. The company wants to provide its customers with different versions of content based on the devices that the customers use to access the website. Which combination of actions should a solutions architect take to meet these requirements ? (Choose

two.)

- A.Configure Amazon CloudFront to cache multiple versions of the content.
- B.Configure a host header in a Network Load Balancer to forward traffic to different instances.
- C.Configure a Lambda@Edge function to send specific objects to users based on the User-Agent header.
- D.Configure AWS Global Accelerator. Forward requests to a Network Load Balancer (NLB) .Configure the NLB to set up host-based routing to different EC2 instances.
- E.Configure AWS Global Accelerator. Forward requests to a Network Load Balancer (NLB) .Configure the NLB to set up path-based routing to different EC2 instances.

Answer : AC

Explanation : B is wrong-NLBs do not understand HTTP (Layer 7/Application layer) headers, this is what ALBs do. Moreover, a host header is information of the DESTINATION server, not the SOURCE client-D and E are wrong-Global Accelerator helps to SPEED UP requests. Doesn't help with CONTENT CUSTOMIZATION.

QUESTION 907

Question : An ecommerce company needs to run a scheduled daily job to aggregate and filter sales records for analytics. The company stores the sales records in an Amazon S3 bucket. Each object can be up to 10 GB in size. Based on the number of sales events, the job can take up to an hour to complete. The CPU and memory usage of the job are constant and are known in advance. A solutions architect needs to minimize the amount of operational effort that is needed for the job to run. Which solution meets these requirements ?

- A.Create an AWS Lambda function that has an Amazon EventBridge (Amazon Cloud Watch Events) notification. Schedule the EventBridge (Cloud Watch Events) event to run once a day.
- B.Create an AWS Lambda function. Create an Amazon API Gateway HTTP API and integrate the API with the function. Create an Amazon EventBridge (Amazon Cloud Watch Events) scheduled event that calls the API and invokes the function.
- C.Create an Amazon Elastic Container Service (Amazon ECS) cluster with an AWS Fargate launch type. Create an Amazon EventBridge (Amazon Cloud Watch Events)

scheduled event that launches an ECS task on the cluster to run the job.

D.Create an Amazon Elastic Container Service (Amazon ECS) cluster with an Amazon EC2 launch type and an Auto Scaling group with at least one EC2 instance.Create an Amazon EventBridge (Amazon Cloud Watch Events) scheduled event that launches an ECS task on the cluster to run the job.

Answer : C

Explanation : -A and B are wrong."work might take up to an hour", hence Lambda (with 15mins execution limit) is wrong-D is wrong."reduce the amount of operational work", hence Fargate (serverless) is preferred over EC2 (self-managed instances) launch type ; C:Correct, less operation and when you have"The job's CPU and memory requirements are consistent and known in advance".You can use Fargate mode

(choose CPU, memory size when set the task definition

QUESTION 908

Question : A company wants to enforce strict security guidelines on accessing AWS Cloudresources as thecompany migrates production workloads from its data centers.Company management wants all users to receive permissions according to their job roles and functions.Which solution meets these requirements with the LEAST operational overhead ?

- A.Create an AWS Single Sign-On deployment.Connect to the on-premises Active Directory to centrally manage users and permissions across the company.
- B.Create an IAM role for each job function.Require each employee to call the sts:AssumeRole action in the AWS Management Console to perform their job role.
- C.Create individual IAM user accounts for each employee.Create an IAM policy for eachjob function, and attach the policy to all IAM users based on their job role.
- D.Create individual IAM user accounts for each employee.Create IAM policies for each job function.Create IAM groups, and attach associated policies to each group.Assign the IAM users to a group based on their job role.

Answer : D

QUESTION 909

Question : A company serves its website by using an Auto Scaling group of Amazon EC2

instances in a single AWS Region. The website does not require a database. The company is expanding, and the company's engineering team deploys the website to a second Region. The company wants to distribute traffic across both Regions to accommodate growth and for disaster recovery purposes. The solution should not serve traffic from a Region in which the website is unhealthy. Which policy or resource should the company use to meet these requirements ?

- A.An Amazon Route 53 simple routing policy.
- B.An Amazon Route 53 multivalue answer routing policy
- C.An Application Load Balancer in one Region with a target group that specifies the EC2 instance IDs from both Regions
- D.An Application Load Balancer in one Region with a target group that specifies the IP addresses of the EC2 instances from both Regions

Answer : B.

Explanation : <https://aws.amazon.com/premiumsupport/knowledge-center/multivalue-versus-simple-policies>

QUESTION 910

Question : A company is running a photo hosting service in the us-east-1 Region. The service enables users across multiple countries to upload and view photos. Some photos are heavily viewed for months, and others are viewed for less than a week. The application allows uploads of up to 20 MB for each photo. The service uses the photo metadata to determine which photos to display to each user. Which solution provides the appropriate user access MOST cost-effectively ?

- A.Store the photos in Amazon DynamoDB. Turn on DynamoDB Accelerator (DAX) to cache frequently viewed items.
- B.Store the photos in the Amazon S3 Intelligent-Tiering storage class. Store the photo metadata and its S3 location in DynamoDB.
- C.Store the photos in the Amazon S3 Standard storage class. Set up an S3 Lifecycle policy to move photos older than 30 days to the S3 Standard-Infrequent Access (S3Standard-IA) storage class. Use the object tags to keep track of metadata.
- D.Store the photos in the Amazon S3 Glacier storage class. Set up an S3 Lifecycle policy to move photos older than 30 days to the S3 Glacier Deep Archive storage class. Store the photo metadata and its S3 location in Amazon Elasticsearch

Service (Amazon ES)

Answer : B

Explanation : S3 Glacier Instant Retrieval-Use for archiving data that is rarely accessed and requires milliseconds retrieval.Data stored in the S3 Glacier Instant Retrieval storage class offers a cost savings compared to the S3 Standard-IA storage class, with the same latency and throughput performance as the S3 Standard-IA storage class.S3 Glacier Instant Retrieval has higher data access costs than S3 Standard-IA.For pricing information, see Amazon S3 pricing.S3 Glacier Instant Retrieval.

Answer : B.Explanation : S3 Glacier Instant Retrieval-Use for archiving data that is rarely accessed and requires milliseconds retrieval.Data stored in the S3 Glacier Instant Retrieval storage class offers a cost savings compared to the S3 Standard-IA storage class, with the same latency and throughput performance as the S3 Standard-IA storage class.S3 Glacier Instant Retrieval has higher data access costs than S3 Standard-IA.For pricing information, see Amazon S3 pricing.S3 Glacier Instant Retrieval.

QUESTION 911

Question : A trucking company is deploying an application that will track the GPS coordinates of all the company's trucks.The company needs a solution that will generate real-time statistics based on metadata lookups with high read throughput and microsecond latency.The database must be fault tolerant and must minimize operational overhead and development effort.Which combination of steps should a solutions architect take to meet these requirements ? (Choose two.)

- A.Use Amazon DynamoDB as the database.
- B.Use Amazon Aurora MySQL as the database.
- C.Use Amazon RDS for MySQL as the database
- D.Use Amazon ElastiCache as the caching layer.
- E.Use Amazon DynamoDB Accelerator (DAX) as the caching layer.

Answer : AE

Explanation : Though DynamoDB offers consistent single-digit-millisecond latency, DynamoDB+DAX takes performance to the next level with response times in microseconds for millions of requests per second for read-heavy workloads.

QUESTION 912

Question : A company is deploying an application that processes large quantities of

data in parallel. The company plans to use Amazon EC2 instances for the workload. The network architecture must be configurable to prevent groups of nodes from sharing the same underlying hardware. Which networking solution meets these requirements ?

- A.Run the EC2 instances in a spread placement group.
- B.Group the EC2 instances in separate accounts.
- C.Configure the EC2 instances with dedicated tenancy.
- D.Configure the EC2 instances with shared tenancy.

Answer : A

QUESTION 913

Question : A solutions architect is designing the cloud architecture for a new application that is being deployed on AWS. The application's users will interactively download and upload files. Files that are more than 90 days old will be accessed less frequently than newer files, but all files need to be instantly available. The solutions architect must ensure that the application can scale to store petabytes of data with maximum durability. Which solution meets these requirements ?

- A.Store the files in Amazon S3 Standard. Create an S3 Lifecycle policy that moves objects that are more than 90 days old to S3 Glacier.
- B.Store the files in Amazon S3 Standard. Create an S3 Lifecycle policy that moves objects that are more than 90 days old to S3 Standard-Infrequent Access (S3 Standard-IA)
- C.Store the files in Amazon Elastic Block Store (Amazon EBS) volumes. Schedule snapshots of the volumes. Use the snapshots to archive data that is more than 90 days old.
- D.Store the files in RAID-striped Amazon Elastic Block Store (Amazon EBS) volumes. Schedule snapshots of the volumes. Use the snapshots to archive data that is more than 90 days old.

Answer : B

QUESTION 914

Question : A company is developing a serverless web application that gives users the ability to interact with real-time analytics from online games. The data from the games must be streamed in real life. The company needs a durable, low-latency database option for user data. The company does not know how many users will use the application. Any design considerations must provide response times of single-digit

milliseconds as the application scales.Which combination of AWS services will meet these requirements ?

(Choose two.)

- A.Amazon CloudFront
- B.Amazon DynamoDB
- C.Amazon Kinesis
- D.Amazon RDS
- E.AWS Global Accelerator

Answer : BC

QUESTION 915

Question : A company is designing an application.The application uses an AWS Lambda function to receive information through Amazon API Gateway and to store the information in an Amazon Aurora PostgreSQL database.During the proof-of-concept stage, the company has to increase the Lambda quotas significantly to handle the high volumes of data that the company needs to load into the database.A solutions architect must recommend a new design to improve scalability and minimize the configuration effort.

Which solution will meet these requirements ?

- A.Refactor the Lambda function code to Apache Tomcat code that runs on Amazon EC2

instances.Connect the database by using native Java Database Connectivity (JDBC) drivers.

- B.Change the platform from Aurora to Amazon DynamoDB.Provision a DynamoDB Accelerator (DAX) cluster.Use the DAX client SDK to point the existing DynamoDB API calls at the DAX cluster

C.Set up two Lambda functions.Configure one function to receive the information.Configure the other function to load the information into the database.Integrate the Lambda functions by using Amazon Simple Notification Service (Amazon SNS) .

- D.Set up two Lambda functions.Configure one function to receive the information.Configure the other function to load the information into the database.Integrate the Lambda functions by using an Amazon Simple Queue Service (Amazon SQS) queue.

Answer : D

QUESTION 916

Question : A company needs to build a reporting solution on AWS. The solution must support SQL queries that data analysts run on the data. The data analysts will run fewer than 10 total queries each day. The company generates 3 GB of new data daily in an on-premises relational database. This data needs to be transferred to AWS to perform reporting tasks. What should a solutions architect recommend to meet these requirements at the LOWEST cost ?

- A.Use AWS Database Migration Service (AWS DMS) to replicate the data from the on-premises database into Amazon S3. Use Amazon Athena to query the data.
- B.Use an Amazon Kinesis Data Firehose delivery stream to deliver the data into an Amazon Elasticsearch Service (Amazon ES) cluster. Run the queries in Amazon ES.
- C.Export a daily copy of the data from the on-premises database. Use an AWS Storage Gateway file gateway to store and copy the export into Amazon S3. Use an Amazon EMR cluster to query the data.
- D.Use AWS Database Migration Service (AWS DMS) to replicate the data from the on-premises database and load it into an Amazon Redshift cluster. Use the Amazon Redshift cluster to query the data.

Answer : D

QUESTION 917

Question : A company is using an Application Load Balancer (ALB) to present its application to the internet. The company finds abnormal traffic access patterns across the application. A solutions architect needs to improve visibility into the infrastructure to help the company understand these abnormalities better. What is the MOST operationally efficient solution that meets these requirements ?

- A.Create a table in Amazon Athena for AWS Cloud Trail logs. Create a query for the relevant information
- B.Enable ALB access logging to Amazon S3. Create a table in Amazon Athena, and query the logs.
- C.Enable ALB access logging to Amazon S3. Open each file in a text editor, and search each line for the relevant information
- D.Use Amazon EMR on a dedicated Amazon EC2 instance to directly query the ALB to acquire traffic access log information.

Answer : B

QUESTION 918

Question : A company wants to monitor its AWS costs for financial review. The cloud operations team is designing an architecture in the AWS Organizations management account to query AWS Cost and Usage Reports for all member accounts. The team must run this query once a month and provide a detailed analysis of the bill. Which solution is the MOST scalable and cost-effective way to meet these requirements ?

- A.Enable Cost and Usage Reports in the management account. Deliver reports to Amazon Kinesis. Use Amazon EMR for analysis.
- B.Enable Cost and Usage Reports in the management account. Deliver the reports to Amazon S3. Use Amazon Athena for analysis.
- C.Enable Cost and Usage Reports for member accounts. Deliver the reports to Amazon S3. Use Amazon Redshift for analysis.
- D.Enable Cost and Usage Reports for member accounts. Deliver the reports to Amazon Kinesis. Use Amazon QuickSight for analysis.

Answer : B

QUESTION 919

Question : A company wants to move from many standalone AWS accounts to a consolidated, multi-account architecture. The company plans to create many new AWS

accounts for different business units. The company needs to authenticate access to these AWS accounts by using a centralized corporate directory service. Which combination of actions should a solutions architect recommend to meet these requirements ? (Choose two.)

- A.Create a new organization in AWS Organizations with all features turned on. Create the new AWS accounts in the organization.
- B.Set up an Amazon Cognito identity pool. Configure AWS Single Sign-On to accept Amazon Cognito authentication.
- C.Configure a service control policy (SCP) to manage the AWS accounts. Add AWS Single Sign-On to AWS Directory Service.
- D.Create a new organization in AWS Organizations. Configure the organization's authentication mechanism to use AWS Directory Service directly.

E.Set up AWS Single Sign-On (AWS SSO) in the organization.Configure AWS SSO, and integrate it with the company's corporate directory service.

Answer : BC

QUESTION 920

Question : A solutions architect needs to design a system to store client case files.The files are core company assets and are important.The number of files will grow over time.The files must be simultaneously accessible from multiple application servers that run on Amazon EC2 instances.The solution must have built-in redundancy.Which solution meets these requirements ?

- A.Amazon Elastic File System (Amazon EFS)
- B.Amazon Elastic Block Store (Amazon EBS)
- C.Amazon S3 Glacier Deep Archive
- D.AWS Backup

Answer : A

QUESTION 921

Question : A solutions architect is designing a shared storage solution for a web application that is deployed across multiple Availability Zones.The web application runs on Amazon EC2 instances that are in an Auto Scaling group.The company plans to make frequent changes to the content.The solution must have strong consistency in returning the new content as soon as the changes occur.Which solutions meet these requirements ? (Choose two.

- A.Use AWS Storage Gateway Volume Gateway Internet Small Computer Systems Interface (iSCSI) block storage that is mounted to the individual EC2 instances.
- B.Create an Amazon Elastic File System (Amazon EFS) file system.Mount the EFS file system on the individual EC2 instances.
- C.Create a shared Amazon Elastic Block Store (Amazon EBS) volume.Mount the EBS volume on the individual EC2 instances.
- D.Use AWS DataSync to perform continuous synchronization of data between EC2 hosts in the Auto Scaling group.

E.Create an Amazon S3 bucket to store the web content.Set the metadata for the Cache-Control header to no-cache.Use Amazon CloudFront to deliver the content.

Answer : AB

QUESTION 922

Question : A company has a legacy data processing application that runs on Amazon EC2 instances.Data is processed sequentially, but the order of results does not matter.The application uses a monolithic architecture.The only way that the company can scale the application to meet increased demand is to increase the size of the instances.The company's developers have decided to rewrite the application to use a microservices architecture on Amazon Elastic Container Service (Amazon ECS) .What should a solutions architect recommend for communication between the microservices ?

A.Create an Amazon Simple Queue Service (Amazon SQS) queue.Add code to the data producers, and send data to the queue.Add code to the data consumers to process data from the queue.

B.Create an Amazon Simple Notification Service (Amazon SNS) topic.Add code to the data producers, and publish notifications to the topic.Add code to the data consumers to subscribe to the topic.

C.Create an AWS Lambda function to pass messages.Add code to the data producers to call the Lambda function with a data object.Add code to the data consumers to receive a data object that is passed from the Lambda function.

D.Create an Amazon DynamoDB table.Enable DynamoDB Streams.Add code to the data producers to insert data into the table.Add code to the data consumers to use the DynamoDB Streams API to detect new table entries and retrieve the data.

Answer : A

QUESTION 923

Question : A large media company hosts a web application on AWS.The company wants to start caching confidential media files so that users around the world will have reliable access to the files.The content is stored in Amazon S3 buckets.The company must deliver the content quickly, regardless of where the requests originate geographically.Which solution will meet these requirements ?

A.Use AWS DataSync to connect the S3 buckets to the web application.

- B.Deploy AWS Global Accelerator to connect the S3 buckets to the web application.
- C.Deploy Amazon CloudFront to connect the S3 buckets to CloudFront edge servers.
- D.Use Amazon Simple Queue Service (Amazon SQS) to connect the S3 buckets to the web application.

Answer : C

QUESTION 924

Question : A solutions architect is designing a workload that will store hourly energy consumption by business tenants in a building. The sensors will feed a database through HTTP requests that will add up usage for each tenant. The solutions architect must use managed services when possible. The workload will receive more features in the future as the solutions architect adds independent components. Which solution will meet these requirements with the LEAST operational overhead ?

- A.Use Amazon API Gateway with AWS Lambda functions to receive the data from the sensors, process the data, and store the data in an Amazon DynamoDB table.
- B.Use an Elastic Load Balancer that is supported by an Auto Scaling group of Amazon EC2 instances to receive and process the data from the sensors. Use an Amazon S3 bucket to store the processed data.
- C.Use Amazon API Gateway with AWS Lambda functions to receive the data from the sensors, process the data, and store the data in a Microsoft SQL Server Express database on an Amazon EC2 instance.
- D.Use an Elastic Load Balancer that is supported by an Auto Scaling group of Amazon EC2 instances to receive and process the data from the sensors. Use an Amazon Elastic File System (Amazon EFS) shared file system to store the processed data.

Answer : A

QUESTION 925

Question : A solutions architect is migrating a document management workload to AWS. The workload keeps 7 TiB of contract documents on a shared storage file system and tracks them on an external database. Most of the documents are stored and retrieved eventually for reference in the future. The application cannot be modified during the migration, and the storage solution must be highly available. Documents are retrieved and stored by web servers that run on Amazon EC2 instances in an Auto Scaling group. The Auto Scaling group can have up to 12 instances. Which solution meets these

requirements MOST cost-effectively ?

- A.Provision an enhanced networking optimized EC2 instance to serve as a shared NFS storage system.
- B.Create an Amazon S3 bucket that uses the S3 Standard-Infrequent Access (S3 Standard-IA) storage class.Mount the S3 bucket to the EC2 instances in the Auto Scaling group.
- C.Create an SFTP server endpoint by using AWS Transfer for SFTP and an Amazon S3 bucket.Configure the EC2 instances in the Auto Scaling group to connect to the SFTP server.
- D.Create an Amazon Elastic File System (Amazon EFS) file system that uses the EFS Standard-Infrequent Access (EFS Standard-IA) storage class.Mount the file system to the EC2 instances in the Auto Scaling group.

Answer : D

QUESTION 926

Question : A company wants to relocate its on-premises MySQL database to AWS.The database accepts regular imports from a client-facing application, which causes a high volume of write operations.The company is concerned that the amount of traffic might be causing performance issues within the application.How should a solutions architect design the architecture on AWS ?

- A.Provision an Amazon RDS for MySQL DB instance with Provisioned IOPS SSD storage.Monitor write operation metrics by using Amazon Cloud Watch.Adjust the provisioned IOPS if necessary.
- B.Provision an Amazon RDS for MySQL DB instance with General Purpose SSD storage.Place anAmazon ElastiCache cluster in front of the DB instance.Configure the application to query ElastiCache instead.
- C.Provision an Amazon DocumentDB (with MongoDB compatibility) instance with a memory optimized instance type.Monitor Amazon Cloud Watch for performance-related issues.Change the instance class if necessary.
- D.Provision an Amazon Elastic File System (Amazon EFS) file system in General Purpose performance mode.Monitor Amazon CloudWatch for IOPS bottlenecks.Change to Provisioned Throughput performance mode if necessary.

Answer : A

QUESTION 927

Question : A company has a remote factory that has unreliable connectivity. The factory needs to gather and process machine data and sensor data so that it can sense products on its conveyor belts and initiate a robotic movement to direct the products to the right location. Predictable low-latency compute processing is essential for the on-premises control systems. Which solution should the factory use to process the data ?

- A.Amazon CloudFront Lambda@Edge functions
- B.An Amazon EC2 instance that has enhanced networking enabled
- C.An Amazon EC2 instance that uses an AWS Global Accelerator
- D.An Amazon Elastic Block Store (Amazon EBS) volume on an AWS Snowball Edge cluster

Answer : A

QUESTION 928

Question : A company is creating a prototype of an ecommerce website on AWS. The website consists of an Application Load Balancer, an Auto Scaling group of Amazon EC2 instances for web servers, and an Amazon RDS for MySQL DB instance that runs with the Single-AZ configuration. The website is slow to respond during searches of the product catalog. The product catalog is a group of tables in the MySQL database that the company does not update frequently. A solutions architect has determined that the CPU utilization on the DB instance is high when product catalog searches occur. What should the solutions architect recommend to improve the performance of the website during searches of the product catalog ?

- A.Migrate the product catalog to an Amazon Redshift database. Use the COPY command to load the product catalog tables.
- B.Implement an Amazon ElastiCache for Redis cluster to cache the product catalog. Use lazy loading to populate the cache.
- C.Add an additional scaling policy to the Auto Scaling group to launch additional EC2 instances when database response is slow.
- D.Turn on the Multi-AZ configuration for the DB instance. Configure the EC2 instances to throttle the product catalog queries that are sent to the database.

Answer : A.

Explanation : https://docs.aws.amazon.com/redshift/latest/dg/r_COPY.html

QUESTION 929

Question : A company wants to use AWS Systems Manager to manage a fleet of Amazon EC2 instances. According to the company's security requirements, no EC2 instances can have internet access. A solutions architect needs to design network connectivity from the EC2 instances to Systems Manager while fulfilling this security obligation. Which solution will meet these requirements ?

- A. Deploy the EC2 instances into a private subnet with no route to the internet.
- B. Configure an interface VPC endpoint for Systems Manager. Update routes to use the endpoint.
- C. Deploy a NAT gateway into a public subnet. Configure private subnets with a default route to the NAT gateway.
- D. Deploy an internet gateway. Configure a network ACL to deny traffic to all destinations except Systems Manager.

Answer : B

Explanation : <https://docs.aws.amazon.com/systems-manager/latest/userguide/setup-create-vpc.html>

QUESTION 930

Question: A company uses AWS to run all components of its three-tier application. The company wants to automatically detect any potential security breaches within the environment. The company wants to track any findings and notify administrators if a potential breach occurs. Which solution meets these requirements ?

- A. Set up AWS WAF to evaluate suspicious web traffic. Create AWS Lambda functions to log any findings in Amazon Cloud Watch and send email notifications to administrators.
- B. Set up AWS Shield to evaluate suspicious web traffic. Create AWS Lambda functions to log any findings in Amazon Cloud Watch and send email notifications to administrators.
- C. Deploy Amazon Inspector to monitor the environment and generate findings in Amazon CloudWatch. Configure an Amazon EventBridge (Amazon Cloud Watch Events) rule to publish a message to an Amazon Simple Notification Service (Amazon SNS) topic to notify administrators by email.
- D. Deploy Amazon GuardDuty to monitor the environment and generate findings in Amazon Cloud Watch. Configure an Amazon EventBridge (Amazon Cloud Watch

Events)

rule to publish a message to an Amazon Simple Notification Service (Amazon SNS) topic to notify administrators by email.

Answer : C

QUESTION 931

Question : A company hosts its enterprise content management platform in one AWS Region but needs to operate the platform across multiple Regions. The company has an Amazon Elastic Kubernetes Service (Amazon EKS) cluster that runs its microservices. The EKS cluster stores and retrieves objects from Amazon S3. The EKS cluster also stores and retrieves metadata from Amazon DynamoDB. Which combination of steps should a solutions architect take to deploy the platform across multiple Regions ? (Choose two.)

- A.Replicate the EKS cluster with cross-Region replication.
- B.Use Amazon API Gateway to create a global endpoint to the EKS cluster.
- C.Use AWS Global Accelerator endpoints to distribute the traffic to multiple Regions.
- D.Use Amazon S3 access points to give access to the objects across multiple Regions. Configure DynamoDB Accelerator (DAX) . Connect DAX to the relevant tables.
- E.Deploy an EKS cluster and an S3 bucket in another Region. Configure cross-Region replication on both S3 buckets. Turn on global tables for DynamoDB.

Answer : BE

QUESTION 932

Question : A company is hosting a web application from an Amazon S3 bucket. The application uses Amazon Cognito as an identity provider to authenticate users and return a JSON Web Token (JWT) that provides access to protected resources that are stored in another S3 bucket. Upon deployment of the application, users report errors and are unable to access the protected content. A solutions architect must resolve this issue by providing proper permissions so that users can access the protected content. Which solution meets these requirements ?

- A.Update the Amazon Cognito identity pool to assume the proper IAM role for access to the protected content.
- B.Update the S3ACI to allow the application to access the protected content.

C.Redeploy the application to Amazon S3 to prevent eventually consistent reads in the S3

bucket from affecting the ability of users to access the protected content.

D.Update the Amazon Cognito pool to use custom attribute mappings within the identity pool and grant users the proper permissions to access the protected content.

Answer : D

QUESTION 933

Question : A company has a new mobile app Anywhere in the world, users can see local news on topics they choose.Users also can post photos and videos from inside the app.

Users access content often in the first minutes after the content is posted.New content quickly replaces older content, and then the older content disappears.The local nature of the news means that users consume 90%of the content within the AWS Region where it is uploaded.Which solution will optimize the user experience by providing the LOWEST latency for content uploads ?

A.Upload and store content in Amazon S3 Use Amazon CloudFront for the uploads.

B.Upload and store content in Amazon S3 Use S3 Transfer Acceleration for the uploads. C.Upload content to Amazon EC2 instances in the Region that is closest to the user.Copy the data to Amazon S3.

D.Upload and store content in Amazon S3 in the Region that is closest to the user.Use multiple distributions of Amazon CloudFront

Answer : A

QUESTION 934

Question : A company runs a containerized application on a Kubernetes cluster in an on-premises data center.The company is using a MongoDB database for data storage.The company wants to migrate some of these environments to AWS, but no code changes or deployment method changes are possible at this time.The company needs a solution that minimizes operational overhead.Which solution meets these requirements ?

A.Use Amazon Elastic Container Service (Amazon ECS) with Amazon EC2 worker nodes for compute and MongoDB on EC2 for data storage.

B.Use Amazon Elastic Container Service (Amazon ECS) with AWS Fargate for

compute and Amazon DynamoDB for data storage.

C.Use Amazon Elastic Kubemnetes Service (Amazon EKS) with Amazon EC2 worker

nodes for compute and Amazon DynamoDB for data storage.

D.Use Amazon Elastic Kubemnetes Service (Amazon EKS) with AWS Fargate for compute and Amazon DocumentDB (with MongoDB compatibility) for data storage.

Answer : D.

Explanation : No changes in the code and operation means they need to use same/similar DB in the AWS so it is Document DB with Mongo DB comp and without having operational diff it means fargate

QUESTION 935

A company is building a new furniture inventory application. The company has deployed the application on a fleet of Amazon EC2 instances across multiple Availability Zones. The EC2 instances run behind an Application Load Balancer (ALB) in their VPC. A solutions architect has observed that incoming traffic seems to favor one EC2 instance resulting in latency for some requests. What should the solutions architect do to resolve this issue ?

A.Disable session affinity (sticky sessions) on the ALB

B.Replace the ALB with a Network Load Balancer

C.increase the number of EC2 instances in each Availability Zone

D.Adjust the frequency of the health checks on the ALB's target group

Answer : B

QUESTION 936

A startup company is using the AWS Cloud to develop a traffic control monitoring system for a large city. The system must be highly available and must provide near-real-time results for residents and city officials even during peak events. Gigabytes of data will come in daily from IoT devices that run at intersections and freeway ramps across the city. The system must process the data sequentially to provide the correct timeline. However results need to show only what has happened in the last 24 hours. Which solution will meet these requirements MOST cost-effectively ?

A.Deploy Amazon Kinesis Data Firehose to accept incoming data from the IoT devices and write the data to Amazon S3. Build a web dashboard to display the data from the

last 24 hours

- B.Deploy an Amazon API Gateway API endpoint and an AWS Lambda function to process incoming data from the IoT devices and store the data in Amazon DynamoDB.Build a web dashboard to display the data from the last 24 hours
- C.Deploy an Amazon API Gateway API endpoint and an Amazon Simple Notification Service (Amazon SNS) topic to process incoming data from the IoT devices Write the data to Amazon Redshift.Build a web dashboard to display the data from the last 24 hours
- D.Deploy an Amazon Simple Queue Service (Amazon SQS) FIFO queue and an AWS Lambda function to process incoming data from the IoT devices and store the data in an Amazon RDS DB instance.Build a web dashboard to display the data from the last 24 hours

Answer : D

QUESTION 937

A company has three AWS accounts Management Development and Production.These accounts use AWS services only in the us-east-1 Region.All accounts have a VPC with VPC Flow Logs configured to publish data to an Amazon S3 bucket in each separate account.For compliance reasons the company needs an ongoing method to aggregate all the VPC flow logs across all accounts into one destination S3 bucket in the Management account.What should a solutions architect do to meet these requirements with the LEAST operational overhead ?

- A.Add S3 Same-Region Replication rules in each S3 bucket that stores VPC flow logs to replicate objects to the destination S3 bucket Configure the destination S3 bucket to allow objects to be received from the S3 buckets in other accounts
- B.Set up an IAM user in the Management account.Grant permissions to the IAM user to access the S3 buckets that contain the VPC flow logs.Run the aws s3 sync command in the AWS CLI to copy the objects to the destination S3 bucket
- C.Use an S3 inventory report to specify which objects in the S3 buckets to copy.Perform an S3 batch operation to copy the objects into the destination S3 bucket in the Management account with a single request
- D.Create an AWS Lambda function in the Management account Grant S3 GET permissions on the source S3 buckets.Grant S3 PUT permissions on the destination S3 bucket.Configure the function to invoke when objects are loaded in the source S3

buckets

Answer : A

QUESTION 938

A company is running a multi-tier web application on AWS. The application runs its database on Amazon Aurora MySQL. The application and database tiers are in the us-east-1 Region. A database administrator who monitors the Aurora DB cluster finds that an intermittent increase in read traffic is creating high CPU utilization on the read replica. The result is increased read latency for the application. The memory and disk utilization of the DB instance are stable throughout the event of increased latency. What should a solutions architect do to improve the read scalability?

- A.Reboot the DB cluster
- B.Create a cross-Region read replica
- C.Configure Aurora Auto Scaling for the read replica
- D.Increase the provisioned read IOPS for the DB instance

Answer : B

QUESTION 939

A developer is creating an AWS Lambda function to perform dynamic updates to a database when an item is added to an Amazon Simple Queue Service (Amazon SQS) queue. A solutions architect must recommend a solution that tracks any usage of database credentials in AWS Cloud Trail. The solution also must provide auditing capabilities. Which solution will meet these requirements?

- A.Store the encrypted credentials in a Lambda environment variable
- B.Create an Amazon DynamoDB table to store the credentials. Encrypt the table.
- C.Store the credentials as a secure string in AWS Systems Manager Parameter Store
- D.Use an AWS Key Management Service (AWS KMS) key store to store the credentials

Answer : D

QUESTION 940

A company has a service that reads and writes large amounts of data from an Amazon S3 bucket in the same AWS Region. The service is deployed on Amazon EC2 instances within the private subnet of a VPC. The service communicates with Amazon S3 over a NAT gateway in the public subnet. However, the company wants a solution that will

reduce the data output costs.Which solution will meet these requirements MOST cost-effectively ?

- A.Provision a dedicated EC2 NAT instance in the public subnet.Configure the route table for the private subnet to use the elastic network interface of this instance as the destination for all S3 traffic
- B.Provision a dedicated EC2 NAT instance in the private subnet.Configure the route table for the public subnet to use the elastic network interface of this instance as the destination for all S3 traffic.
- C.Provision a VPC gateway endpoint.Configure the route table for the private subnet to use the gateway endpoint as the route for all S3 traffic.
- D.Provision a second NAT gateway.Configure the route table for the private subnet to use this NAT gateway as the destination for all S3 traffic.

Answer : C

QUESTION 941

A company has an application that uses an Amazon DynamoDB table for storage.A solutions architect discovers that many requests to the table are not returning the latest data.The company's users have not reported any other issues with database performance Latency is in an acceptable range.Which design change should the solutions architect recommend ?

- A.Add read replicas to the table.
- B.Use a global secondary index (GSI) .
- C.Request strongly consistent reads for the table
- D.Request eventually consistent reads for the table.

Answer : C

QUESTION 942

A company has developed a new content-sharing application that runs on Amazon Elastic Container Service (Amazon ECS) .The application runs on Amazon Linux Docker tasks that use the Amazon EC2 launch type.The application requires a storage solution that has the following characteristics:-Accessibility or multiple ECS tasks through bind mounts ; -Resiliency across Availability Zones, -Burstable throughput of up to 3 Gbps, -Ability to be scaled up over time Which storage solution meets these requirements ?

- A.Launch an Amazon FSx for Windows File Server Multi-AZ instance.Configure the ECS task definitions to mount the Amazon FSx instance volume at launch
- B.Launch an Amazon Elastic File System (Amazon EFS) instance.Configure the ECS task definitions to mount the EFS Instance volume at launch.
- C.Create a Provisioned IOPS SSD (io2) Amazon Elastic Block Store (Amazon EBS) volume with Multi-Attach set to enabled.Attach the EBS volume to the ECS EC2 instance Configure ECS task definitions to mount the EBS instance volume at launch.
- D.Launch an EC2 instance with several Provisioned IOPS SSD (io2) Amazon Elastic Block Store (Amazon EBS) volumes attached m a RAID O configuration.Configure the EC2 instance as an NFS storage server.Configure ECS task definitions to mount the volumes at launch.

Answer : B

QUESTION 943

An airline that is based in the United States provides services for routes in North America and Europe.The airline is developing a new read-intensive application that customers can use to find flights on either continent.The application requires strong read consistency and needs scalable database capacity to accommodate changes in user demand.The airline needs the database service to synchronize with the least possible latency between the two continents and to provide a simple failover mechanism to a second AWS Region.Which solution will meet these requirements ?

- A.Deploy Microsoft SQL Server on Amazon EC2 instances in a Region in North America.Use SOL Server binary log replication on an EC2 instance in a Region in Europe.
- B.Create an Amazon DynamoDB global table Add a Region from North America and a Region from Europe to the table.Query data with strongly consistent reads.
- C.Use an Amazon Aurora MySQL global database.Deploy the read-write node in a Region in North America, and deploy read-only endpoints in Regions in North America and Europe.Query data with global read consistency.
- D.Create a subscriber application that uses Amazon Kinesis Data Streams for an Amazon Redshift cluster in a Region in North America.Create asecond subscriber application for the Amazon Redshift cluster in a Region in Europe.Process all database modifications through Kinesis Data Streams.

Answer : C

QUESTION 944

A disaster relief company is designing a new solution to analyze real-time csv data. The data is collected by a network of thousands of research stations met are distributed across the world. The data volume is consistent and constant, and the size of each data We is 512 KB. The company needs to stream the data and analyze the data in real time. Which combination of actions should a solutions architect take to meet these requirements ? (Select TWO.)

- A.Provision an appropriately sized Amazon Simple Queue Service (Amazon SOS) queue. Use the AWS SDK at the research stations to write the data into the SOS queue
- B.Provision an appropriately sized Amazon Kinesis Data Firehose delivery stream. Use the AWS SDK at the research stations to write the data into the delivery stream and then into an Amazon S3 bucket.
- C.Provision an appropriately sized Amazon Kinesis Data Analytics application. Use the AWS CLI to configure Kinesis Data Analytics with SOL queries
- D.Provision an AWS Lambda function to process the data. Set up the BatchSize property on the Lambda event source.
- E.Provision an AWS Lambda function to process the data. Set up an Amazon EventBridge (Amazon CloudWatch Events) cron expression rule to invoke the Lambda function

Answer : AD

QUESTION 945

A company stores can wordings on a monthly basis. Users access lie recorded files randomly within 1 year of recording, but users rarely access the files after 1 year. The company wants to optimize its solution by allowing only files that ant newer than 1 year old to be queried and retrieved as quickly as possible. A delay in retrieving older fees is acceptable. Which solution meets these requirements MOST cost-effectively ?

- A.Store individual files in Amazon S3 Glacier. Store search metadata in object tags that are created in S3 Glacier. Query the S3 Glacier tags to retrieve the files from S3 Glacier.
- B.Store individual files in Amazon S3. Use S3 Lifecycle polices to move the ties to S3 Glacier after 1 year. Query and retrieve the files that are in Amazon S3 by using Amazon Athena. Query and retrieve the files that are in S3 Glacier by using S3 Glacier Select.
- C.Store Individual files In Amazon \$3. Store search metadata for each archive In

Amazon S3.Use S3 Lifecycle policies to move the ties to S3 Glacier after 1 year.Query and retrieve tie flies by searching for metadata from Amazon S3.

D.Store individual files in Amazon S3 Use S3.Lifecycle policies to move the files to S3 Glacier after 1 year.Store search metadata in Amazon RDS Query the Sea from Amazon RDS Retrieve the files from Amazon S3 or S3 Glacier.

Answer : D

QUESTION 946

A company needs to save the results from a medical trial to an Amazon S3 repository.The repository must allow a few scientists to add new dies and must restrict all other users to read-only access No users can have the ability to modify or delete any files in the repository.The company must heap every lie in the repository for a minimum of 1 year after its creation date.Which solution will meet these requirements ?

- A.Use S3 Object Lock In governance mode with a legal hold of 1 year
- B.Use S3 Object Lock in compliance mode with a retention period of 365 days.
- C.Use an IAM role to restrict all users from deleting or changing objects in the S3 bucket.Use an S3 bucket policy to only allow the IAM role
- D.Configure the S3 bucket to invoke an AWS Lambda function every tune an object is added.Configure the function to track the hash of the saved object to that modified objects can be marked accordingly

Answer : B

QUESTION 947

A company runs an internet-facing web application on AWS.The company uses Amazon Route 53 for DNS management and has a public hosted zone lo route traffic from the internet to the application.The company wants to tog DNS response codes to help system administrators perform any root cause analysis in the future.Which solution will meet these requirements ?

- A.Use Route 53 to configure query toggling
- B.Use AWS Cloud Trail lo record ail Route 53 queries
- C.Use Amazon Cloud Watch to record and process Route 53 metrics
- D.Use AWS Trusted Advisor to perform on-demand root cause analysis

Answer : A

QUESTION 948

Some of the company's customers are retrieving records frequently, leading to an increase in costs for the company. The company wants to limit retrieved requests in the future. The company also wants to ensure that if one customer reaches its retrieval limit other customers will not be affected. Which solution will meet these requirements?

- A. Set up server-side throttling limits for API Gateway.
- B. Limit DynamoDB read throughput on the table to an amount that results in the maximum cost that the company is willing to incur.
- C. Set up a usage plan for API Gateway. Implement throttling limits for each customer and distribute API keys to each customer.
- D. Set up AWS Budgets. Monitor the usage of API Gateway and DynamoDB. Configure an alarm to provide an alert when the cost exceeds a certain threshold each month.

Answer : D

QUESTION 949

A company is preparing to store confidential data in Amazon S3. For compliance reasons the data must be encrypted at rest. Encryption key usage must be logged for auditing purposes. Keys must be rotated every year. Which solution meets these requirements and is the MOST operationally efficient?

- A. Server-side encryption with customer-provided keys (SSE-C)
- B. Server-side encryption with Amazon S3 managed keys (SSE-S3)
- C. Server-side encryption with AWS KMS (SSE-KMS) customer master keys (CMKs) with manual rotation
- D. Server-side encryption with AWS KMS (SSE-KMS) customer master keys (CMKs) with automatic rotation

Answer : D

QUESTION 950

A company has a corporate network on premises and has three VPCs in the AWS Cloud. The company has one VPC each for development, test, and production. The company wants its system administrators to securely gain remote command-line access from the corporate network to Amazon EC2 instances in the VPCs. Which solution meets these requirements MOST cost-effectively?

- A.Set up a VPN connection between the corporate network and each of the three VPCs by using AWS VPN.Use Remote Desktop Protocol (RDP) or SSH over the VPN connection to access the EC2 instances remotely.
- B.Configure the EC2 instances to use an instance profile that trusts AWS Systems Manager.Use Systems Manager Session Manager to gain console access to the EC2 instances.
- C.Create a new VPC Purchase and install a virtual router from AWS Marketplace. Establish a VPN connection from the corporate network to this router.Establish another VPN connection from the router to the other three VPCs.Use Remote Desktop Protocol (RDP) or SSH over the VPN connection to access the EC2 instances remotely.
- D.Create a new VPC Establish a VPN connection to the new VPC.Configure peering connections between the new VPC and the existing VPCs.Inthe new VPC create an EC2bastion host to serve as a jump box lo EC2 instances in the other VPCs.Use Remote Desktop Protocol (RDP) or SSH over the VPN connection to the bastion host

Answer : A

QUESTION 951

A company uses a simple static website and wants to host it on AWS.The company already has a domain that it uses for email.The company needs a hosting solution that supports HTTPS.Which solution will meet these requirements MOST cost-effectively ?

A.Create an Amazon S3 bucket with a name to match the website.Upload the website to the S3 bucket.Set up website hosting for the S3 bucket.Set up the DNS to point to the S3 website endpoint.

B.Create an Amazon S3 bucket upload the website to the S3 bucket.Set up an HTTPS certificate by using AWS Certificate Manager (ACM) .Create in Amazon CloudFront distribution for the S3 bucket and choose Price Class All.

C.Set up an open-source content management system (CMS) from AWS Marketplace Deploy the CMS across two Availability Zones Copy the website onto the CMS Set up the DNS to point to me CMS

D.Create an Amazon S3 bucket Upload the website to the S3 bucket Set up an HTTPS certificate by using AWS Certificate Manager (ACM) Create an Amazon CloudFront distribution for the S3 bucket and choose Price Class 100 Point to the CloudFront distribution

Answer : D

QUESTION 952

A solutions architect must secure a VPC network that hosts Amazon EC2 instances. The EC2 instances contain highly sensitive data and run in a private subnet. According to company policy the EC2 instances may run in the VPC can access only approved third-party software repositories on the internet for software product updates that use the third party's URL. Other internet traffic must be blocked. Which solution meets these requirements?

- A.Update the route table for the private subnet to route the outbound traffic to an AWS Network Firewall. Configure domain list rule groups
- B.Set up an AWS WAF web ACL. Create a custom set of rules that filter traffic requests based on source and destination IP address range sets.
- C.Implement strict inbound security group roles. Configure an outbound rule that allows traffic only to the authorized software repositories on the internet by specifying the URLs
- D.Configure an Application Load Balancer (ALB) in front of the EC2 instances. Direct an outbound traffic to the ALB. Use a URL-based rule listener in the ALB's target group for outbound access to the internet

Answer : C

QUESTION 953

A company has a data ingestion workflow that includes the following components:-An Amazon Simple Notification Service (Amazon SNS) topic that receives notifications about new data deliveries-An AWS Lambda function that processes and stores the data. The ingestion workflow occasionally fails because of network connectivity issues. When failure occurs the corresponding data is not ingested unless the company manually reruns the job. What should a solutions architect do to ensure that all notifications are eventually processed?

- A.Configure the Lambda function for deployment across multiple Availability Zones
- B.Modify the Lambda function's configuration to increase the CPU and memory allocations for the function.
- C.Configure the SNS topic's retry strategy to increase both the number of retries and

the wait time between retries

D.Configure an Amazon Simple Queue Service (Amazon SQS) queue as the on failure destination. Modify the Lambda function to process messages in the queue.

Answer : A

QUESTION 954

A company runs a highly available image-processing application on Amazon EC2 instances in a single VPC. The EC2 instances run inside several subnets across multiple Availability Zones. The EC2 instances do not communicate with each other. However, the EC2 instances download images from Amazon S3 and upload images to Amazon S3 through a single NAT gateway. The company is concerned about data transfer charges. What is the MOST cost-effective way for the company to avoid Regional data transfer charges?

- A.Launch the NAT gateway in each Availability Zone
- B.Replace the NAT gateway with a NAT instance
- C.Deploy a gateway VPC endpoint for Amazon S3
- D.Provision an EC2 Dedicated Host to run the EC2 instances

Answer : A

QUESTION 955

A company wants to migrate an on-premises data center to AWS. The data center hosts an SFTP server that stores its data on an NFS-based file system. The server holds 200 GB of data that needs to be transferred. The server must be hosted on an Amazon EC2 instance that uses an Amazon Elastic File System (Amazon EFS) file system. Which combination of steps should a solutions architect take to automate this task? (Select TWO)

- A.Launch the EC2 instance into the same Availability Zone as the EFS file system
- B.Install an AWS DataSync agent in the on-premises data center
- C.Create a secondary Amazon Elastic Block Store (Amazon EBS) volume on the EC2 instance for the data
- D.Manually use an operating system copy command to push the data to the EC2 instance
- E.Use AWS DataSync to create a suitable location configuration for the on-premises SFTP server

Answer : CD

QUESTION 956

A company hosts a multi-tier web application on Amazon Linux Amazon EC2 instances behind an Application Load Balancer. The instances run in an Auto Scaling group across multiple Availability Zones. The company observes that the Auto Scaling group launches more On-Demand Instances when the application's end users access high volumes of static web content. The company wants to optimize cost. What should a solutions architect do to redesign the application MOST cost-effectively ?

- A. Update the Auto Scaling group to use Reserved Instances instead of On-Demand Instances
- B. Update the Auto Scaling group to scale by launching Spot Instances instead of On-Demand Instances
- C. Create an Amazon CloudFront distribution to host the static web contents from an Amazon S3 bucket
- D. Create an AWS Lambda function behind an Amazon API Gateway API to host the static website contents

Answer : B

QUESTION 957

A company runs a website that uses a content management system (CMS) on Amazon EC2. The CMS runs on a single EC2 instance and uses an Amazon Aurora MySQL Multi-AZ DB instance for the data for Website images are stored on an Amazon Elastic Block Store (Amazon EBS) volume that is mounted inside the EC2 instance. Which combination of actions should a solutions architect take to improve the performance and resilience of the website ? (Select TWO)

- A. Move the website images into an Amazon S3 bucket that is mounted on every EC2 instance
- B. Share the website images by using an NFS share from the primary EC2 instance Mount this share on the other EC2 instances
- C. Move the website images onto an Amazon Elastic File System (Amazon EFS) file system that is mounted on every EC2 instance
- D. Create an Amazon Machine Image (AMI) from the existing EC2 instance Use the AMI to provision new instances behind an Application Load Balancer as part of an Auto Scaling group Configure the Auto Scaling group to maintain a minimum of two instances Configure an accelerator in AWS Global Accelerator for the website

E.Create an Amazon Machine Image (AMI) from the existing EC2 instance.Use the AMI to provision new instances behind an Application Load Balancer as part of an Auto Scaling group Configure the Auto Scaling group to maintain a minimum of two instances Configure an Amazon CloudFront distribution for the website

Answer : CE

QUESTION 958

A company wants to migrate its accounting system from an on-premises data center to the AWS Cloud in a single AWS Region.Data security and an immutable audit log are the top priorities.The company must monitor all AWS activities for compliance auditing.The company has enabled AWS Cloud Trail but wants to make sure it meets these requirements.Which actions should a solutions architect take to protect and secure Cloud Trail ? (Choose two.

- A.Enable Cloud Trail log file validation.
- B.Install the Cloud Trail Processing Library.
- C.Enable logging of Insights events in Cloud Trail.
- D.Enable custom logging from the on-premises resources.
- E.Create an AWS Config rule to monitor whether Cloud Trail is configured to use server-side encryption with AWS KMS managed encryption keys (SSE-KMS) .

Answer : C E

QUESTION 959

A company has an AWS Lambda function that needs read access to an Amazon S3 bucket that is located in the same AWS account.Which solution will meet, these requirements in the MOST secure manner ?

- A.Apply an S3 bucket policy that grants read access to the S3 bucket.
- B.Apply an IAM role to the Lambda function.Apply an IAM policy to the role to grant read access to the S3 bucket
- C.Embed an access key and a secret key in the Lambda function's code to grant the required IAM permissions for read access to the S3 bucket.
- D.Apply an IAM role to the Lambda function.Apply an IAM policy to the role to grant read access to all S3 buckets in the account.

Answer : D

Explanation : <https://aws.amazon.com/premiumsupport/knowledge-center/access->

[denied-lambda-s3-b ucket/](#)

QUESTION 960

A company runs its infrastructure on AWS and has a registered base of 700,000 users for its document management application. The company intends to create a product that converts large .pdf files to .jpg image files. The .pdf files average 5 MB in size. The company needs to store the original files and the converted files. A solutions architect must design a scalable solution to accommodate demand that will grow rapidly over time. Which solution meets these requirements MOST cost-effectively ?

- A. Save the .pdf files to Amazon S3. Configure an S3 PUT event to invoke an AWS Lambda function to convert the files to .jpg format and store them back in Amazon S3.
- B. Save the pdf files to Amazon DynamoDB. Use the DynamoDB Streams feature to invoke an AWS Lambda function to convert the files to jpg format and store them back in DynamoDB.
- C. Upload the .pdf files to an AWS Elastic Beanstalk application that includes Amazon EC2 instances, Amazon Elastic Block Store (Amazon EBS) storage, and an Auto Scaling group. Use a program in the EC2 instances to convert the files to .jpg format. Save the pdf files and the jpg files in the EBS store.
- D. Upload the pdf files to an AWS Elastic Beanstalk application that includes Amazon EC2 instances, Amazon Elastic File System (Amazon EFS) storage, and an Auto Scaling group. Use a program in the EC2 instances to convert the file to .jpg format. Save the pdf files and the jpg files in the EFS store.

Answer : B

QUESTION 961

A company is designing a new multi-tier web application that consists of the following components:- Web and application servers that run on Amazon EC2 instances as part of Auto Scaling groups - An Amazon RDS DB instance for data storage A solutions architect needs to limit access to the application servers so that only the web servers can access them. Which solution will meet these requirements ?

- A. Deploy AWS PrivateLink in front of the application servers. Configure the network ACL to allow only the web servers to access the application servers
- B. Deploy a VPC endpoint in front of the application servers. Configure the security group to allow only the web servers to access the application servers

C.Deploy a Network Load Balancer with a target group that contains the application servers"Auto Scaling group.Configure the networkACI to allow only the web servers to access the application servers

D.Deploy an Application Load Balancer with a target group that contains the application servers'Auto Scaling group Configure the security group to allow only the web servers to access the application servers.

Answer : D

QUESTION 962

A company is planning to store sensitive documents in an Amazon S3 bucket.The documents must be encrypted al rest.The company wants to manage the underlying keys that are used for encryption.However, the company does not want to manage the encryption and decryption process.Which solutions will meet these requirements ? (Select TWO.)

- A.Use server-side encryption with customer-provided encryption keys (SSE-C) .
- B.Use client-side encryption with AWS managed keys.
- C.Use server-side encryption with S3 managed encryption keys (SSE-S3)
- D.Use server-side encryption with AWS KMS managed encryption keys (SSE-KMS) with a key policy document that is 40 KB in size
- E.Use server-side encryption with AWS KMS managed encryption keys (SSE-KMS) that the company uploads to AWS KMS.

Answer : CE

QUESTION 963

A company has a popular gaming platform running on AWS.The application is sensitive to latency because latency can impact the user experience and introduce unfair advantages to some players.The application is deployed in a very AWS Region.It runs on AmazonFC2 Instances which are part of Auto Scaling groups configured behind Application Load Balancers (ALBs) .A solutions architect needs to implement a mechanism to monitor the health of the application and redirect traffic to healthy endpoints.Which solution meets these requirements ?

- A.Configure an accelerator In AWS Global Accelerator Add a listener for the port that the application listens on and attach it to a Regional endpoint in each Region Add the ALB as the endpoint

B.Create an Amazon CloudFront distribution and specify the ALB as the origin server
Configure the cache behaviour to use origin cache headers Use AWS Lambda functions to optimize the traffic

C.Create an Amazon CloudFront distribution and specify Amazon S3 as the origin server.Configure tie cache behaviour to use origin cache headers Use AWS Lambda functions to optimize the traffic

D.Configure an Amazon DynamoDB database to serve as the data store for the application Create a DynamoDB Accelerator (DAX) cluster to act as the in-memory cache for DynamoDB hosting the application data

Answer : A

QUESTION 964

A company has NFS servers in an on-premises data center that need to periodically back up small amounts of data to Amazon S3.Which solution meets these requirement and is MOST cost-effective ?

- A.Set up AWS Glue to copy the data from the on-premises servers to Amazon S3.
- B.Set up an AWS DataSync agent on the on-premises servers, and sync the data to Amazon S3
- C.Set up an SFTP sync using AWS Transfer for SFTP to sync data from on premises to Amazon S3
- D.Set up an AWS Direct Connect connection between the on-premises data center and a VPC, and copy the data to Amazon S3

Answer : B

QUESTION 965

A company uses an application to present metrics from sporting events to the public.The application must scale quickly during live events and must store these metrics for log-term reporting purposes.The company's architecture includes the following:-Amazon EC2 instances that run in an Auto Scaling group in private subnets-A network Load Balancer That runs in public subnets-A MongoDB database cluster that runs across multiple EC2 instances A solutions architect must implement a solution that minimizes operational overhead.The solution also must be able to fail automatically.What should the solutions architect set up to meet these requirements ?

- A.An Amazon DynamoDB database

B.An Amazon RDS for MySQL D6 instance

C.EC2 instances that run MySQL

D.Amazon Redshift

Answer : A

QUESTION 966

A company is running a mission-critical application on Amazon EC2 instances behind an Application Load Balancer. The instances run in an Auto Scaling group in a single AWS Region. The application is using a database in Amazon Aurora as the data tier. A recent audit revealed that the current deployment of Aurora is not highly available. What should a solutions architect do to improve the availability of the database ?

A.Configure an Aurora Replica

B.Configure storage replication.

C.Configure storage auto scaling.

D.Configure cross-Region replication

Answer : A

QUESTION 967

A company runs a stateless web application in production on a group of Amazon EC2 On-Demand Instances behind an Application Load Balancer. The application experiences heavy usage during an 8-hour period each business day. Application usage is moderate and steady overnight. Application usage is low during weekends. The company wants to minimize its EC2 costs without affecting the availability of the application. Which solution will meet these requirements ?

A.Use Spot Instances for the entire workload.

B.Use Reserved instances for the baseline level of usage. Use Spot Instances for any additional capacity that the application needs.

C.Use On-Demand Instances for the baseline level of usage. Use Spot Instances for any additional capacity that the application needs

D.Use Dedicated Instances for the baseline level of usage. Use On-Demand Instances for any additional capacity that the application needs

Answer : B

QUESTION 968

An ecommerce company uses an Amazon Aurora DB cluster to store customer transactions. The company also maintains a separate Amazon DynamoDB table that contains item sales information. The company wants the DB cluster to invoke a recently deployed AWS Lambda function to update the DynamoDB table every time a row is inserted into the database. Which combination of steps should a solutions architect take to meet these requirements ? (Select TWO.)

- A.Modify the Lambda function to allow outbound communication to the DB cluster
- B.Modify the DB cluster to allow outbound communication to the Lambda function.
- C.Modify the DB cluster to allow outbound communication to the DynamoDB table.
- D.Ensure that the DB cluster has an IAM role that allows the DB cluster to invoke Lambda functions.
- E.Ensure that the Lambda function has an IAM role that allows Lambda to invoke functions on the DB cluster.

Answer : CE

QUESTION 969

A solutions architect creates a VPC that includes two public subnets and two private subnets. A corporate security mandate requires the solutions architect to launch all Amazon EC2 instances in a private subnet. However when the solutions architect launches an EC2 instance that runs a web server on ports 80 and 443 in a private subnet, no external internet traffic can connect to the server. What should the solutions architect do to resolve this issue ?

- A.Attach the EC2 instance to an Auto Scaling group in a private subnet Ensure that the DNS record for the website resolves to the Auto Scaling group identifier
- B.Provision an internet-facing Application Load Balancer (ALB) in a public subnet Add the EC2 instance to the target group that is associated with the ALB Ensure that the DNS record for the website resolves to the ALB
- C.Launch a NAT gateway in a private subnet Update the route table for the private subnets to add a default route to the NAT gateway Attach a public Elastic IP address to the NAT gateway
- D.Ensure that the security group that is attached to the EC2 instance allows HTTP traffic on port 80 and HTTPS traffic on port 443 Ensure that the DNS record for the website resolves to the public IP address of the EC2 instance.

Answer : A

QUESTION 970

A company is experiencing sudden increases in demand. The company needs to provision large Amazon EC2 instances from an Amazon Machine image (AMI). The instances will run in an Auto Scaling group. The company needs a solution that provides minimum initialization latency to meet the demand. Which solution meets these requirements?

- A. Use the aws ec2 register-image command to create an AMI from a snapshot. Use AWS Step Functions to replace the AMI in the Auto Scaling group.
- B. Enable Amazon Elastic Block Store (Amazon EBS) fast snapshot restore on a snapshot. Provision an AMI by using the snapshot. Replace the AMI in the Auto Scaling group with the new AMI.
- C. Enable AMI creation and define lifecycle rules in Amazon Data Lifecycle Manager (Amazon DLM). Create an AWS Lambda function that modifies the AMI in the Auto Scaling group.
- D. Use Amazon EventBridge (Amazon Cloud Watch Events) to invoke AWS Backup lifecycle policies that provision AMIs. Configure Auto Scaling group capacity limits as an event source in EventBridge (CloudWatch Events).

Answer : A

QUESTION 971

A company hosts a marketing website in an on-premises data center. The website consists of static documents and runs on a single server. An administrator updates the website content infrequently and uses an SFTP client to upload new documents. The company decides to host its website on AWS and to use Amazon CloudFront. The company's solutions architect creates a CloudFront distribution. The solutions architect must design the most cost-effective and resilient architecture for website hosting to serve as the CloudFront origin. Which solution will meet these requirements?

- A. Create a virtual server by using Amazon Lightsail. Configure the web server in the Lightsail instance. Upload website content by using an SFTP client.
- B. Create an AWS Auto Scaling group for Amazon EC2 instances. Use an Application Load Balancer. Upload website content by using an SFTP client.
- C. Create a private Amazon S3 bucket. Use an S3 bucket policy to allow access from a CloudFront origin access identity (OAI). Upload website content by using the AWS

CLI

D.Create a public Amazon S3 bucket
Configure AWS Transfer for SFTP
Configure the S3 bucket for website hosting
Upload website content by using the SFTP client

Answer : C

QUESTION 972

A pharmaceutical company is developing a new drug. The volume of data that the company generates has grown exponentially over the past few months. The company's researchers regularly require a subset of the entire dataset to be immediately available with minimal lag. However, the entire dataset does not need to be accessed on a daily basis. All the data currently resides in on-premises storage arrays, and the company wants to reduce ongoing capital expenses. Which storage solution should a solutions architect recommend to meet these requirements ?

- A.Run AWS DataSync as a scheduled cron job to migrate the data to an Amazon S3 bucket on an ongoing basis
- B.Deploy an AWS Storage Gateway file gateway with an Amazon S3 bucket as the target storage. Migrate the data to the Storage Gateway appliance.
- C.Deploy an AWS Storage Gateway volume gateway with cached volumes with an Amazon S3 bucket as the target storage. Migrate the data to the Storage Gateway appliance.
- D.Configure an AWS Site-to-Site VPN connection from the on-premises environment to AWS. Migrate data to an Amazon Elastic File System (Amazon EFS) file system.

Answer : B

QUESTION 973

A company manages and runs a critical data management application in containers that are hosted on Amazon Elastic Container Service (Amazon ECS) . The application has endpoints that are exposed through Application Load Balancers(ALBs) . The application uses an Amazon Elastic File System (Amazon EFS) file system mount for persistent data storage. The company has configured Amazon ECS to use a minimal IAM instance role. Which combination of actions should a solutions architect take to improve the overall security posture of the application ? (Select TWO.)

- A.Decompose the Amazon ECS IAM instance role. Use only ECS task roles.
- B.Enable EFS encryption in transit to protect data that is being written to Amazon EFS.

- C.Use AWS Config to define patch management policies on the container instances.
- D.Use Amazon Macie integration with Amazon EFS to monitor and protect sensitive information in the file system.
- E.Use Amazon GuardDuty to authenticate data access between the ALBs and the container instances.

Answer : CD

QUESTION 974

A company is upgrading its critical web-based application. The application is hosted on Amazon EC2 instances that are part of an Auto Scaling group behind an Application Load Balancer (ALB). The company wants to test the new configurations with a specific amount of traffic before the company begins to route all traffic to the upgraded application. How should a solutions architect design the architecture to meet these requirements?

- A.Create a new launch template. Associate the new launch template with the Auto Scaling group. Attach the Auto Scaling group to the ALB. Distribute traffic by using redirect rules.
- B.Create a new launch template. Create an additional Auto Scaling group. Associate the new launch template with the additional Auto Scaling group. Attach the additional Auto Scaling group to the ALB. Distribute traffic by using weighted target groups.
- C.Create a new launch template. Create an additional Auto Scaling group. Associate the new launch template with the additional Auto Scaling group. Create an additional ALB. Attach the additional Auto Scaling group to the additional ALB. Use an Amazon Route 53 failover routing policy to route traffic.
- D.Create a new launch template. Create an additional Auto Scaling group. Associate the new launch template with the additional Auto Scaling group. Create an additional ALB. Attach the additional Auto Scaling group to the additional ALB. Use an Amazon Route 53 weighted routing policy to route traffic.

Answer : A

QUESTION 975

A company runs an online ticketing application with backend services that run on Amazon EC2 instances. The EC2 instances belong to an Auto Scaling group and run behind an Application Load Balancer. The application experiences periods of high user

traffic when a popular event is posted online. The company wants a solution that will be able to handle increases in user traffic without affecting the user experience. What should a solutions architect do to meet these requirements ?

- A.Configure a scheduled scaling policy for peak hours with a recurrence schedule set to every day.
- B.Configure a target tracking scaling policy that uses the average aggregate CPU utilization target metric.
- C.Configure a step scaling policy that is based on an Amazon CloudWatch alarm that monitors CPU utilization.
- D.Configure an Application Load Balancer health check that increases the Auto Scaling group capacity whenever the application returns HTTP 503 error codes.

Answer : B

QUESTION 976

A company runs a shopping application that uses Amazon DynamoDB to store customer information. In case of data corruption, a solutions architect needs to design a solution that meets a recovery point objective (RPO) of 15 minutes and a recovery time objective (RTO) of 1 hour. What should the solutions architect recommend to meet these requirements ?

- A.Configure DynamoDB global tables. For RPO recovery, point the application to a different AWS Region.
- B.Configure DynamoDB point-in-time recovery. For RPO recovery, restore to the desired point in time.
- C.Export the DynamoDB data to Amazon S3 Glacier on a daily basis. For RPO recovery, import the data from S3 Glacier to DynamoDB.
- D.Schedule Amazon Elastic Block Store (Amazon EBS) snapshots for the DynamoDB table every 15 minutes. For RPO recovery, restore the DynamoDB table by using the EBS snapshot.

Answer : B

QUESTION 977

A company runs demonstration environments for its customers on Amazon EC2 instances. Each environment is isolated in its own VPC. The company's operations team needs to be notified when RDP or SSH access to an environment has been

established.What should a solutions architect recommend to meet these requirements ?

A.Configure Amazon CloudWatch Application Insights to create AWS Systems Manager OpsItems when RDP or SSH access is detected.

B.Configure the EC2 instances with an IAM instance profile that has an IAM role with the AmazonSSMManagedInstanceCore policy attached.

C.Publish VPC flow logs to Amazon CloudWatch Logs.Create required metric filters.Create an Amazon CloudWatch metric alarm with a notification action for when the alarm is in the ALARM state.

D.Configure an Amazon EventBridge (Amazon Cloud Watch Events) rule to listen for events of type EC2 Instance State-change Notification.Configure an Amazon Simple Notification Service (Amazon SNS) topic as a target.Subscribe the operations team to the topic.

Answer : C

QUESTION 978

A company is planning to move a dynamic web application to AWS.Application assets are stored in a Linux file server.The total volume of data is 140TB, and the company has 100 Mbps of internet bandwidth available.The company does not want to make any changes to the application during migration.File permissions must be preserved.How should a solutions architect migrate and store the data to meet these requirements ?

A.Transfer the data by using the file interface for AWS Snowball.Use Amazon Elastic File System (Amazon EFS) as the migration destination and for storage.

B.Transfer the data by using the Amazon S3 interface for AWS Snowball.Use Amazon FSx for Lustre as the migration destination and for storage.

C.Transfer the data by using the AWS CLI s3 cp command.Use Amazon EC2 with Amazon Elastic Block Store (Amazon EBS) as the migration destination and for storage.

D.Transfer the data by using AWS DataSync.Use AWS Storage Gateway File Gateway as the migration destination and for storage.

Answer : A

QUESTION 979

A company wants its public web application to run on Amazon EC2 instances in an Auto Scaling group behind an Application Load Balancer (ALB) .The application

must use a publicly trusted SSL certificate.Which solution will meet these requirements MOST cost-effectively ?

- A.Provision a public SSL/TLS certificate through AWS Certificate Manager (ACM) .Configure the new certificate on the HTTPS listener for the ALB.
- B.Use AWS Certificate Manager Private Certificate Authority to issue an SSL/TLS certificate.Configure the new certificate on the HTTPS listener for the ALB.
- C.Create a self-signed certificate on one of the EC2 instances in the Auto Scaling group.Export the certificate, and configure it on the HTTPS listener for the ALB.
- D.Deploy an EC2-hosted certificate authority (CA) .Import a trusted root certificate.Issue a new SSL/TLS certificate.Configure the new certificate on the HTTPS listener for the ALB.

Answer : A

QUESTION 980

A company needs to store data from its healthcare application.The application's data frequently changes.A new regulation requires audit access at all levels of the stored data.The company hosts the application on an on-premises infrastructure that is running out of storage capacity.A solutions architect must securely migrate the existing data to AWS while satisfying the new regulation.Which solution will meet these requirements ?

- A.Use AWS DataSync to move the existing data to Amazon S3.Use AWS Cloud Trail to log data events.
- B.Use AWS Snowcone to move the existing data to Amazon S3.Use AWS Cloud Trail to log management events.
- C.Use Amazon S3 Transfer Acceleration to move the existing data to Amazon S3.Use AWS Cloud Trail to log data events.
- D.Use AWS Storage Gateway to move the existing data to Amazon S3.Use AWS Cloud Trail to log management events.

Answer : A

QUESTION 981

A large international university has deployed all of its compute services in the AWS Cloud.These services include Amazon EC2,Amazon RDS, and Amazon DynamoDB.The university currently relies on many custom scripts to back up its

infrastructure. However, the university wants to centralize management and automate data backups as much as possible by using AWS native options. Which solution will meet these requirements ?

- A. Use third-party backup software with an AWS Storage Gateway tape gateway virtual tape library.
- B. Use AWS Backup to configure and monitor all backups for the services in use.
- C. Use AWS Config to set lifecycle management to take snapshots of all data sources on a schedule.
- D. Use AWS Systems Manager State Manager to manage the configuration and monitoring of backup tasks.

Answer : B

QUESTION 982

A company is hosting a static website on Amazon S3 and is using Amazon Route 53 for DNS. The website is experiencing increased demand from around the world. The company must decrease latency for users who access the website. Which solution meets these requirements MOST cost-effectively ?

- A. Replicate the S3 bucket that contains the website to all AWS Regions. Add Route 53 geolocation routing entries.
- B. Provision accelerators in AWS Global Accelerator. Associate the supplied IP addresses with the S3 bucket. Edit the Route 53 entries to point to the IP addresses of the accelerators.
- C. Add an Amazon CloudFront distribution in front of the S3 bucket. Edit the Route 53 entries to point to the CloudFront distribution.
- D. Enable S3 Transfer Acceleration on the bucket. Edit the Route 53 entries to point to the new endpoint.

Answer : C

QUESTION 983

A survey company has gathered data for several years from areas the United States. The company hosts the data in an Amazon S3 bucket that is 3 TB in size and growing. The company has started to share the data with a European marketing firm that has S3 buckets. The company wants to ensure that its data transfer costs remain as low as possible. Which solution will meet these requirements ?

- A.Configure the Requester Pays feature on the company's S3 bucket
- B.Configure S3 Cross-Region Replication from the company's S3 bucket to one of the marketing firm's S3 buckets.
- C.Configure cross-account access for the marketing firm so that the marketing firm has access to the company's S3 bucket.
- D.Configure the company's S3 bucket to use S3 Intelligent-Tiering Sync the S3 bucket to one of the marketing firm's S3 buckets

Answer : A

QUESTION 984

A company must save all the email messages that its employees send to customers for a period of 12 months. The messages are stored in a binary format and vary in size from 1 KB to 20 KB. The company has selected Amazon S3 as the storage service for the messages. Which combination of steps will meet these requirements MOST cost-effectively? (Select TWO.)

- A.Create an S3 bucket policy that denies the S3.Delete Object action.
- B.Create an S3 lifecycle configuration that deletes the messages after 12 months.
- C.Upload the messages to Amazon S3. Use S3 Object Lock in governance mode
- D.Upload the messages to Amazon S3. Use S3 Object Lock in compliance mode.
- E.Use S3 Inventory. Create an AWS Batch job that periodically scans the inventory and deletes the messages after 12 months

Answer : BD

QUESTION 985

A company recently migrated multiple applications and databases from an on-premises data center to the AWS Cloud. Most of the applications run on AWS Fargate, and some of the applications run on Amazon EC2 instances. Most of the databases run on Amazon RDS, and a small number of databases run on EC2 Instances. All the applications and databases must be available 24 hours a day, 7 days a week. The company uses AWS Organizations to manage AWS accounts. A solutions architect must recommend how to minimize the cost of these workloads over the next 3 years. Which solution meets these requirements?

- A.Purchase All Upfront Reserved Instances with a 3-year term for Amazon EC2 and Fargate

- B.Purchase All Upfront Reserved Instances with a 3-year term for Amazon EC2 and Amazon RDS
- C.Purchase All Upfront Compute Savings Plans with a 3-year term for Amazon EC2 and Fargate Purchase All Upfront Reserved Instances with a 3-year term for Amazon RDS
- D.Purchase All Upfront EC2 Instance Savings Plans with a 3-year term for Amazon EC2 and Fargate Purchase All Upfront Reserved Instances with a 3-year term for Amazon RDS

Answer : D

QUESTION 986

A company has a data ingestion workflow that consists the following:-An Amazon Simple Notification Service (Amazon SNS) topic for notifications about new data deliveries-An AWS Lambda function to process the data and record metadata. The company observes that the ingestion workflow fails occasionally because of network connectivity issues. When such a failure occurs, the Lambda function does not ingest the corresponding data unless the company manually reruns the job. Which combination of actions should a solutions architect take to ensure that the Lambda function ingests all data in the future ? (Select TWO.)

- A.Deploy the Lambda function In multiple Availability Zones.
- B.Create an Amazon Simple Queue Service (Amazon SQS) queue, and subscribe It to me SNS topic.
- C.Increase the CPU and memory that are allocated to the Lambda function.
- D.Increase provisioned throughput for the Lambda function.
- E.Modify the Lambda function to read from an Amazon Simple Queue Service (Amazon SQS) queue

Answer : BE

QUESTION 987

A company processes images into thumbnails and returns an email confirmation to the end user upon completion. The company's existing solution is facing performance bottlenecks and scalability issues. The company wants to migrate this process to AWS and implement a solution that requires the least possible configuration. Which solution meets these requirements ?

- A.Use Amazon S3 to store images and send notifications to AWS Lambda Configure an AWS Lambda function to process the images into thumbnails, store the thumbnails in Amazon S3, and send an email confirmation through Amazon Simple Email Service (Amazon SES)
- B.Use Amazon S3 to store images and send notifications to Amazon Simple Queue Service (Amazon SQS) Configure an Amazon EC2 instance to poll the SQS queue to process the images into thumbnails, store the thumbnails in Amazon S3, and send an email confirmation through Amazon Simple Email Service (Amazon SES)
- C.Use Amazon S3 to store images and send notifications to Amazon Simple Notification Service (Amazon SNS) Configure Amazon SNS to invoke an AWS Lambda function to process the images into thumbnails, store the thumbnails in Amazon S3, and send an email confirmation through Amazon Simple Email Service (Amazon SES) .
- D.Use Amazon S3 to store images and send notifications to Amazon Simple Queue Service (Amazon SQS) Configure an AWS Lambda function to retrieve the messages from the SQS queue process the images into thumbnails, store the thumbnails in Amazon S3, and send an email confirmation through Amazon Simple Email Service (Amazon SES)

Answer : D

QUESTION 988

A company is hosting its website on Amazon S3 and is using Amazon CloudFront to cache content. The company has an upcoming product launch. An employee accidentally published marketing content to the website before the official release of the product. The company needs to remove the marketing content from the website as quickly as possible. Which solution will meet these requirements ?

- A.Deploy the updated version of the website to another S3 bucket Update the origin for CloudFront
- B.Delete the marketing content in the existing S3 bucket Invalidate the file path in CloudFront
- C.Create a new CloudFront cache policy with a low TTL Associate the new policy with the existing CloudFront distribution
- D.Delete the marketing content in the existing S3 bucket Update the S3 bucket policy to block requests to the file path

Answer : D

QUESTION 989

A company is running a media application in an on-premises data center and has accumulated 500 TB of data. The company needs to migrate the data from the application's existing network-attached file system to AWS. Users rarely access data that is older than 1 year. Which solution meets these requirements MOST cost-effectively?

- A. Use AWS Snowmobile to move the data to Amazon S3 Create an S3 Lifecycle policy to transition data that is older than 1 year to S3 Glacier
- B. Use multiple AWS Snowball Edge Storage Optimized devices to move the data to Amazon S3 Create an S3 Lifecycle policy to transition data that is older than 1 year to S3 Standard-Infrequent Access (S3 Standard-IA)
- C. Set up an AWS Direct Connect connection between the on-premises data center and AWS Transfer the data directly to Amazon S3 by using the Direct Connect connection Create an S3 Lifecycle policy to transition data that is older than 1 year to S3 Glacier
- D. Set up an AWS Site-to-Site VPN connection between the on-premises data center and AWS Transfer the data directly to Amazon S3 by using the Site-to-Site VPN connection Create an S3 Lifecycle policy to transition data that is older than 1 year to S3 Standard-Infrequent Access (S3 Standard-IA)

Answer : B

QUESTION 990

A company needs to migrate a large amount of data from an on-premises storage area network (SAN) to Amazon S3. The SAN currently has 200 TB of data and is receiving an additional 20 TB of data each month. The company has a 500 Mbps connection to the internet. What should the company do to migrate the data to Amazon S3 in the LEAST amount of time?

- A. Use a file syncing application to sync the data to Amazon S3 over the internet through a public S3 endpoint Sync any changed data the same way until the SAN is decommissioned
- B. Use an AWS Snowball Edge Storage Optimized device to migrate the initial 200 TB of data to Amazon S3 Sync any changed data by using AWS DataSync until the SAN

is decommissioned

C.Set up an AWS Site-to-Site VPN connection Use a file syncing application to sync the data to Amazon S3 through a private S3 endpoint Sync any changed data the same way until the SAN is decommissioned

D.Set up a 10 Gbps AWS Direct Connect connection Migrate the initial 200 TB of data to Amazon S3 by using a file syncing application Sync any changed data the same way until the SAN is decommissioned.

Answer : B

QUESTION 991

Which AWS service or feature can be used to find availability status information on all AWS services ?

A.AWS Personal Health

B.Dashboard AWS Cloud Trail

C.AWS Service Health Dashboard

D.Amazon CloudWatch

Answer : A

QUESTION 992

A company wants to design its cloud architecture so that its workloads are resilient, can consistently perform their intended functions correctly, and can recover from failure quickly.Which pillar of the AWS Well-Architected Framework does this architecture represent ?

A.Security

B.Performance efficiency

C.Operational excellence

D.Reliability

Answer : C

QUESTION 993

A company with multiple accounts and teams wants to set up a new multi-account AWS environment.Which AWS service supports this requirement ?

A.AWS CloudFormation

B.AWS Control Tower

C.AWS Config

D.Amazon Virtual Private Cloud (Amazon VPC)

Answer : B

QUESTION 994

A company wants to track its daily AWS resource usage to avoid reaching service quotas unexpectedly. The company needs to receive notifications when any service quota is exceeded. Which combination of actions should a solutions architect take to meet this requirement? (Select TWO.)

A.Configure Amazon Simple Notification Service (Amazon SNS) as the target to send notifications

B.Use the Describe TrustedAdvisorChecks API operation to get AWS Trusted Advisor Service Limits checks every 24 hours

C.Create an AWS Lambda function that runs every 24 hours and refreshes the AWS Trusted Advisor Service Limits checks

D.Use AWS Config to monitor the AWS resources service quotas and create a periodic invocation for an AWS Lambda function.

E.Use Amazon EventBridge (Amazon Cloud Watch Events) to capture the events. Configure Amazon Simple Notification Service (Amazon SNS) as the target

Answer : BD

QUESTION 995

A solutions architect needs to connect a company's corporate network to its VPC to allow on-premises access to its AWS resources. The solution must provide encryption of all traffic between the corporate network and the VPC at the network layer and the session layer. The solution also must provide security controls to prevent unrestricted access between AWS and the on-premises systems. Which solution meets these requirements?

A.Configure AWS Direct Connect to connect to the VPC. Configure the VPC route tables to allow and deny traffic between AWS and on premises as required

B.Create an IAM policy to allow access to the AWS Management Console only from a defined set of corporate IP addresses. Restrict user access based on job responsibility by using an IAM policy and roles.

C.Configure AWS Site-to-Site VPN to connect to the VPC. Configure route table entries

to direct traffic from on premises to the VPC.Configure instance security groups and network ACIs to allow only required traffic from on premises.

D.Configure AWS Transit Gateway to connect to the VPC Configure route table entries to direct traffic from on premises to the VPC.Configure instance security groups and network ACLs to allow only required traffic from on premises.

Answer : A

QUESTION 996

A company asks a solutions architect to design a new network architecture for its on-premises applications.The company has five application VPCs Each application VPC supports a different application.All the application VPCs need to reach a shared services VPC for enterprise services.The application VPCs need to communicate with each other.The number of VPCs will increase as the company grows.The company needs a solution that minimizes operational overhead during setup and maintenance.Which solution will meet these requirements ?

A.Use VPC endpoint services from the application VPCs to reach services on the shared services VPC.

B.Use VPC peering connections between the application VPCs and the shared servicesVPC.Add a peering connection and routes between the application VPCs as needed.

C.Use AWS Transit Gateway to connect the shared services VPC to each of the application VPCs.Configure routing to allow the application VPCs to access the shared services network.

D.Use a virtual private gateway and VPNs in each application VPC to reach the company network.Iransfer the network traffic between the application VPCs by using dynamic routing to the on-premises network.

Answer : C

QUESTION 997

A company has migrated several applications to AWS in the past 3 months.The company wants to know the breakdown of costs for each of these applications.The company wants to receive a regular report that includes this information.Which solution will meet these requirements MOST cost-effectively ?

A.Use AWS Budgets to download data for the past 3 months into a csv file.Look up the

desired information.

- B.Load AWS Cost and Usage Reports into an Amazon RDS DB instance.Run SQL queries to get the desired information.
- C.Tag all the AWS resources with a key for cost and a value of the application's name.Activate cost allocation tags.Use Cost Explorer to get the desired information.
- D.Tag all the AWS resources with a key for cost and a value of the application's name.Use the AWS Billing and Cost Management console to download bills for the past 3 months.Look up the desired information.

Answer : C

QUESTION 998

A company is building a new data analysis application that will ingest large volumes of data into an Amazon S3 bucket.The company is concerned that sensitive information, such as personally identifiable information (PII) .Might be included in some of the data that is ingested.The company needs a solution that will scan for sensitive data and log the findings.What should a solutions architect recommend to meet these requirements ?

- A.Deploy Amazon Inspector to scan the ingested data Configure Amazon Inspector to log findings to Amazon Cloud Watch if Amazon Inspector finds any sensitive data.
- B.Deploy Amazon QuickSight to scan the ingested data.Configure QuickSight to log findings to Amazon CloudWatch if QuickSight finds any sensitive data.
- C.Create a series of AWS Lambda functions to call Amazon GuardDuty to perform scans of the ingested data.If GuardDuty finds any sensitive data, invoke a Lambda function to write findings to Amazon Cloud Watch.
- D.Create a series of AWS Lambda functions to call Amazon Macie to perform scans of the ingested data.If Macie finds any sensitive data, invoke a Lambda function to write findings to Amazon CloudWatch.

Answer : D

QUESTION 999

A company is migrating its application to an Amazon Elastic Kubernetes Service (Amazon EKS) cluster behind an Application Load Balancer (ALB) .The disaster recovery (DR) requirements for the application include the ability to fail over to another AWSRegion with minimal downtime.Which combination of actions should

a solutions architect take to meet this requirement ? (Select TWO.)

- A.Create a scaled-down clone environment in the DR Region.Use auto scaling policies with the EKS nodes.
- B.Create an Amazon Route 53 record that points to the ALB.Configure an active-passive failover routing policy on the record.
- C.Create an AWS Resource Access Manager policy that grants the application users access to the DR environment when the DR environment is needed.
- D.Create an AWS Lambda function that monitors the availability of the main environment and deploys the DR environment when the DR environment is needed.
- E.Create an AWS CloudFormation template that deploys the stack.Deploy the same template in the DR Region when the main environment is unavailable.

Answer : AB

QUESTION 1000

A company is building a shopping application on AWS.The application offers a catalog that changes once each month and needs to scale with traffic volume.The company wants the lowest possible latency from the application.Data from each user's shopping cart needs to be highly available.User session data must be available even if the user is disconnected and reconnects.What should a solutions architect do to ensure that the shopping cart data is preserved at all times ?

- A.Configure an Application Load Balancer to enable the sticky sessions feature (session affinity) for access to the catalog in Amazon Aurora.
- B.Configure Amazon ElastiCache for Redis to cache catalog data from Amazon DynamoDB and shopping cart data from the user's session.
- C.Configure Amazon Elasticsearch Service (Amazon ES) to cache catalog data from Amazon DynamoDB and shopping can data from the user's session.
- D.Configure an Amazon EC2 instance with Amazon Elastic Block Store (Amazon EBS) storage for the catalog and shopping cart.Configure automated snapshots.

Answer : B

QUESTION 1001

A company has an internet-facing application that runs on premises.The application contains mostly user-generated content.Thedata is stored in an on-premises network-attached storage system.The company wants to archive this data annually and has

chosen to move the archival data to Amazon S3. The company needs a solution to migrate the archival data into an S3 bucket. Which solution will meet these requirements?

- A. Use AWS Storage Gateway Volume Gateway. Cache the data, and then replicate the data from the on-premises environment to Amazon S3.
- B. Use AWS DataSync. Create a configuration to replicate the data from the on-premises environment to Amazon S3.
- C. Use AWS Transfer Family. Use an SFTP client to serially transfer the data from the on-premises environment to Amazon S3.
- D. Use Amazon S3 Transfer Acceleration. Use a third-party backup utility to replicate the data from the on-premises environment to Amazon S3.

Answer : B

QUESTION 1002

A company created and hosts a legacy software application for its customers. The application runs on a dedicated Linux server for each customer. The application stores no persistent data except for MySQL data. The company experienced some data corruption issues in the past and wants to move the application to AWS. The company needs to implement a solution to optimize the stability of the application. The solution also must give the company the ability to restore a customer's database to a specific point in time. The company will migrate customer data by using AWS Database Migration Service (AWS DMS). Which architecture should a solutions architect recommend to meet these requirements?

- A. Set up a shared Amazon Aurora database. Configure an Amazon EC2 launch template for each customer.
- B. Set up a shared Amazon Aurora database. Create an Amazon EC2 Amazon Machine Image (AMI) for each customer. Use the AMI to launch the application.
- C. Set up an Amazon RDS database and an Amazon EC2 instance for each customer. Download the installation script. Run the script to install and configure the application.
- D. Set up an Amazon RDS database for each customer. Deploy the application by using an Amazon EC2 launch template. Use user data to configure the customer-specific data.

Answer : C

QUESTION 1003

A company has a web application that users access from around the world. The company has web servers in multiple AWS Regions to support the traffic. A solutions architect must configure an Amazon Route 53 routing policy to send traffic to only the active web servers. Which configuration meets this requirement?

- A.Create a simple routing policy that uses health checks for each Region
- B.Create a multivalue answer routing policy that uses health checks for each Region
- C.Create a geoproximity routing policy with a health check bias of 99 for each Region
- D.Create a weighted routing policy with a health check weight of 100 for each Region

Answer : B

QUESTION 1004

A company has an application workflow that uses an AWS Lambda function to download and decrypt files from Amazon S3. These files are encrypted using AWS Key Management Service Customer Master Keys (AWS KMS CMKs). A solutions architect needs to design a solution that will ensure the required permissions are set correctly. Which combination of actions accomplish this? (Select TWO.)

- A.Attach the kms:decrypt permission to the Lambda function's resource policy.
- B.Grant the decrypt permission for the Lambda IAM role in the KMS key's policy.
- C.Grant the decrypt permission for the Lambda resource policy in the KMS key's policy.
- D.Create a new IAM policy with the kms:decrypt permission and attach the policy to the Lambda function.
- E.Create a new IAM role with the kms:decrypt permission and attach the execution role to the Lambda function.

Answer : BE

QUESTION 1005

A company is planning to host its compute-intensive applications on Amazon EC2 instances. The majority of the network traffic will be between these applications. The company needs a solution that minimizes latency and maximizes network throughput. The underlying hardware for the EC2 instances must not be shared with any other company. Which solution will meet these requirements?

- A.Launch EC2 instances as Dedicated Hosts in a cluster placement group
- B.Launch EC2 instances as Dedicated Hosts in a partition placement group

C.Launch EC2 instances as Dedicated Instances in a cluster placement group

D.Launch EC2 instances as Dedicated Instances in a partition placement group

Answer : A

QUESTION 1006

A company hosts an application on an AWS Lambda function that runs a number of processing steps. The Lambda function typically takes less than 5 minutes to run unless errors occur. The company needs to decouple the application code because of past errors that caused the whole process to fail when a processing step took longer than expected. The processing steps must be available to be replayed up to 12 months from when the original processing occurred. How should a solutions architect design the new solution?

A.Configure Amazon EventBridge (Amazon Cloud Watch Events) , and create an archive Split the processes into separate Lambda functions Create rules for the different event patterns from the Lambda functions to perform processing

B.Keep the Lambda function in place, but increase the timeout to 15 minutes Configure the Lambda function to write each processing step into an Amazon DynamoDB table Replay the steps by using a separate Lambda function and by querying the table when necessary

C.Keep the Lambda function in place, but increase the timeout to 60 minutes. Configure the Lambda function to write each processing step into a daily file in an Amazon S3 bucket Replay the steps by using a separate Lambda function and by querying the file based on required date

D.Configure Amazon Simple Queue Service (Amazon SQS) queues, and create an archive Split the processes into separate Lambda functions Pass messages to different queues as each process is completed, and invoke the next Lambda function to poll the queue for new messages Replay the messages from the SQL queue archive when necessary

Answer : D

QUESTION 1007

A company hosts a three-tier web application that includes a PostgreSQL database. The database stores the metadata from documents. The company searches the metadata for key terms to retrieve documents that the company reviews in a report each month. The

documents are stored in Amazon S3. The documents are usually written only once, but they are updated frequently. The reporting process takes a few hours with the use of relational queries. The reporting process must not affect any document modifications or the addition of new documents. What are the MOST operationally efficient solutions that meet these requirements? (Select TWO)

- A. Set up a new Amazon DocumentDB (with MongoDB compatibility) cluster that includes a read replica. Scale the read replica to generate the reports.
- B. Set up a new Amazon RDS for PostgreSQL Reserved Instance and an On-Demand read replica. Scale the read replica to generate the reports.
- C. Set up a new Amazon Aurora PostgreSQL DB cluster that includes a Reserved Instance and an Aurora Replica. Issue queries to the Aurora Replica to generate the reports.
- D. Set up a new Amazon RDS for PostgreSQL Multi-AZ Reserved Instance. Configure the reporting module to query the secondary RDS node so that the reporting module does not affect the primary node.
- E. Set up a new Amazon DynamoDB table to store the documents. Use a fixed write capacity to support new document entries. Automatically scale the read capacity to support the reports.

Answer : AC

QUESTION 1008

A company needs a backup strategy for its three-tier stateless web application. The web application runs on Amazon EC2 instances in an Auto Scaling group with a dynamic scaling policy that is configured to respond to scaling events. The database tier runs on Amazon RDS for PostgreSQL. The web application does not require temporary local storage on the EC2 instances. The company's recovery point objective (RPO) is 2 hours. The backup strategy must maximize scalability and optimize resource utilization for this environment. Which solution will meet these requirements?

- A. Take snapshots of Amazon Elastic Block Store (Amazon EBS) volumes of the EC2 instances and database every 2 hours to meet the RPO.
- B. Configure a snapshot lifecycle policy to take Amazon Elastic Block Store (Amazon EBS) snapshots. Enable automated backups in Amazon RDS to meet the RPO.
- C. Retain the latest Amazon Machine Images (AMIs) of the web and application tiers. Enable automated backups in Amazon RDS and use point-in-time recovery to meet

the RPO

D.Take snapshots of Amazon Elastic Block Store (Amazon EBS) volumes of the EC2 instances every 2 hours
Enable automated backups in Amazon RDS and use point-in-time recovery to meet the RPO

Answer : D

QUESTION 1009

A company is adopting serverless architecture. The company's solutions architect wants to modernize an application that has source data in csv format. A large team of developers needs to use the application to run SQL queries and reports on demand by joining data across multiple tables. Which combination of actions will meet these requirements MOST cost-effectively ? (Select TWO)

- A.Store the source data in Amazon S3
- B.Load the source data into Amazon RDS
- C.Run on-demand reports and queries by using Amazon Athena
- D.Run on-demand reports and queries by using Amazon QuickSight
- E.Run on-demand reports and queries by using Amazon DynamoDB

Answer : AC

QUESTION 1010

A company allows users to upload and store photos through its website. The website has users from all around the world. All images that users upload are stored in a centralized Amazon S3 bucket. The company wants to increase the speed in which its entire user base can upload photos through the website. What should a solutions architect recommend to meet these requirements ?

- A.Create an Amazon CloudFront distribution. Use the Amazon S3 Standard storage class to store files
- B.Create an Amazon CloudFront distribution. Configure the distribution settings and origin.
- C.Configure S3 Transfer Acceleration on the S3 bucket. Use the standard S3 endpoint to upload files
- D.Configure S3 Transfer Acceleration on the S3 bucket. Use the S3 Accelerate endpoint to upload files

Answer : C

QUESTION 1011

A company runs an on-premises application. The company is planning to move the application to containers by using Kubernetes. The company wants to migrate the application to AWS to reduce the overhead of container infrastructure management. A solutions architect must configure the environment to allow deployment of the company's own custom Amazon Machine Image (AMI) to nodes. Which solution will meet these requirements with the LEAST operational overhead?

- A. Provision a Kubernetes cluster on Amazon EC2 Store container images in Docker Hub
- B. Use AWS Fargate on Amazon Elastic Kubernetes Service (Amazon EKS) Store container images in Amazon Elastic Container Registry (Amazon ECR).
- C. Use Amazon Elastic Kubernetes Service (Amazon EKS) with managed worker nodes Store container images in Amazon Elastic Container Registry (Amazon ECR)
- D. Use Amazon Elastic Kubernetes Service (Amazon EKS) with self-managed worker nodes Store container images in an image repository that runs on Amazon EC2

Answer : A

QUESTION 1012

A company copies 200 TB of data from a recent ocean survey onto AWS Snowball Edge Storage Optimized devices. The company has a high performance computing (HPC) cluster that is hosted on AWS to look for oil and gas deposits. A solutions architect must provide the cluster with consistent sub-millisecond latency and high-throughput access to the data on the Snowball Edge Storage Optimized devices. The company is sending the devices back to AWS. Which solution will meet these requirements?

- A. Create an Amazon S3 bucket Import the data into the S3 bucket Configure an AWS Storage Gateway file gateway to use the S3 bucket Access the file gateway from the HPC cluster instances
- B. Create an Amazon S3 bucket Import the data into the S3 bucket Configure an Amazon FSx for Lustre file system and integrate it with the S3 bucket Access the FSx for Lustre file system from the HPC cluster instances
- C. Create an Amazon S3 bucket and an Amazon Elastic File System (Amazon EFS)

file system Import the data into the S3 bucket Copy the data from the S3 bucket to the EFS file system Access the EFS file system from the HPC cluster instances

D.Create an Amazon FSx for Lustre file system Import the data directly into the FSx for Lustre file system Access the FSx for Lustre file system from the HPC cluster instances

Answer : B

QUESTION 1013

A financial services company wants to shut down two data centers and migrate more than 100 TB of data to AWS. The data has an intricate directory structure with millions of small files stored in deep hierarchies of subfolders. Most of the data is unstructured and the company's file storage consists of SMB-based storage types from multiple vendors. The company does not want to change its applications to access the data after migration. What should a solutions architect do to meet these requirements with the LEAST operational overhead ?

- A.Use AWS Direct Connect to migrate the data to Amazon S3.
- B.Use AWS DataSync to migrate the data to Amazon FSx for Lustre
- C.Use AWS DataSync to migrate the data to Amazon FSx for Windows File Server
- D.Use AWS Direct Connect to migrate the data on-premises file storage to an AWS Storage Gateway volume gateway.

Answer : D

QUESTION 1014

A solutions architect at a company is designing the architecture for a two-tiered web application. The web application is composed of an internet-facing Application Load Balancer (ALB) that forwards traffic to an Auto Scaling group of Amazon EC2 instances. The EC2 instances must be able to access a database that runs on Amazon RDS. The company has requested a defense-in-depth approach to the network layout. The company does not want to rely solely on security groups or network ACLs. Only the minimum resources that are necessary should be routable from the internet. Which network design should the solutions architect recommend to meet these requirements ?

- A.Place the ALB, EC2 instances, and RDS database in private subnets.
- B.Place the ALB in public subnets. Place the EC2 instances and RDS database in private

subnets

C.Place the ALB and EC2 instances in public subnets.Place the RDS database in private subnets.

D.Place the ALB outside the VPC.Place the EC2 instances and RDS database in private subnets.

Answer : B

Explanation : <https://aws.amazon.com/premiumsupport/knowledge-center/public-load-balancer-private-ec2/>

QUESTION 1015

A company needs to send large amounts of data from its data center to an Amazon S3 bucket on a regular basis.The data must be encrypted and must be transferred over a network that provides consistent bandwidth and low latency.What should a solutions architect do to meet these requirements ?

- A.Use an AWS Direct Connect connection
- B.Use an AWS VPN CloudHub connection
- C.Use HTTPS TLS for encryption of data in transit
- D.Use a gateway VPC endpoint to access Amazon S3

Answer : B

QUESTION 1016

A company is planning to deploy a business-critical application in the AWS Cloud.The application requires durable storage with consistent, low-latency performance.Which type of storage should a solutions architect recommend to meet these requirements ?

- A.Instance store volume
- B.Amazon ElastiCache for Memcached cluster
- C.Provisioned IOPS SSD Amazon Elastic Block Store (Amazon EBS) Volume
- D.Throughput Optimized HDD Amazon Elastic Block Store (Amazon EBS) volume

Answer : C

QUESTION 1017

A company has implemented a self-managed DNS solution on three Amazon EC2 instances behind a Network Load Balancer (NLB) in the us-west-2 Region Most of

the company's users are located in the United States and Europe. The company wants to improve the performance and availability of the solution by using an AWS Region in Europe. The company launches and configures three EC2 instances in the eu-west-1 Region and adds the EC2 instances as targets for a new NLB. Which solutions will allow traffic to be routed to all the EC2 instances ? (Select TWO)

- A.Create an Amazon Route 53 geolocation routing policy to route requests to one of the two NLBs Create an Amazon CloudFront distribution Use the Route 53 record as the distribution's origin
- B.Create a standard accelerator by using AWS Global Accelerator Create endpoint groups in us-west-2 and eu-west-1 Add the two NLBs as endpoints for the endpoint groups
- C.Attach Elastic IP addresses to the six EC2 instances Create an Amazon Route 53 geolocation routing policy to route requests to one of the six EC2 instances Create an Amazon CloudFront distribution Use the Route 53 record as the distribution's origin.
- D.Create a standard accelerator by using AWS Global Accelerator Create endpoint groups in us-west-2 and eu-west-1 Add the six EC2 instances directly as endpoints for the endpoint groups Delete the NLBs
- E.Replace the two NLBs with two Application Load Balancers (ALBs) Create an Amazon Route 53 latency routing policy to route requests to one of the two ALBs Create an Amazon CloudFront distribution Use the Route 53 record as the distribution's origin

Answer : AB

QUESTION 1018

An application that is hosted on Amazon EC2 instances needs to access an Amazon S3 bucket. Traffic must not traverse the internet. How should a solutions architect configure access to meet these requirements ?

- A.Create a private hosted zone by using Amazon Route 53
- B.Set up a gateway VPC endpoint for Amazon S3 in the VPC
- C.Configure the EC2 instances to use a NAT gateway to access the S3 bucket
- D.Establish an AWS Site-to-Site VPN connection between the VPC and the S3 bucket

Answer : B

QUESTION 1019

A company experienced a breach that affected several applications in its on-premises data center. The attacker took advantage of vulnerabilities in the custom applications that were running on the servers. The company is now migrating its applications to run on Amazon EC2 instances. The company wants to implement a solution that actively scans for vulnerabilities on the EC2 instances and sends a report that details the findings. Which solution will meet these requirements?

- A. Deploy AWS Shield to scan the EC2 instances for vulnerabilities. Create an AWS Lambda function to log any findings to AWS Cloud Trail.
- B. Deploy Amazon Macie and AWS Lambda functions to scan the EC2 instances for vulnerabilities. Log any findings to AWS Cloud Trail.
- C. Turn on Amazon GuardDuty. Deploy the GuardDuty agents to the EC2 instances. Configure an AWS Lambda function to automate the generation and distribution of reports that detail the findings.
- D. Turn on Amazon Inspector. Deploy the Amazon Inspector agent to the EC2 instances. Configure an AWS Lambda function to automate the generation and distribution of reports that detail the findings.

Answer : D

QUESTION 1021

An ecommerce company is building a distributed application that involves several serverless functions and AWS services to complete order-processing tasks. These tasks require manual approvals as part of the workflow. A solutions architect needs to design an architecture for the order-processing application. The solution must be able to combine multiple AWS Lambda functions into responsive serverless applications. The solution also must orchestrate data and services that run on Amazon EC2 instances, containers, or on-premises servers. Which solution will meet these requirements with the LEAST operational overhead?

- A. Use AWS Step Functions to build the application.
- B. Integrate all the application components in an AWS Glue job.
- C. Use Amazon Simple Queue Service (Amazon SQS) to build the application.
- D. Use AWS Lambda functions and Amazon EventBridge (Amazon CloudWatch Events) events to build the application.

Answer : B

QUESTION 1021

A company runs multiple Amazon EC2 Linux instances in a VPC across two Availability Zones. The instances host applications that use a hierarchical directory structure. The applications need to read and write rapidly and concurrently to shared storage. What should a solutions architect do to meet these requirements?

- A.Create an Amazon S3 bucket Allow access from all the EC2 instances in the VPC
- B.Create an Amazon Elastic File System (Amazon EFS) file system Mount the EFS file system from each EC2 instance
- C.Create a file system on a Provisioned IOPS SSD (io2) Amazon Elastic Block Store (Amazon EBS) volume Attach the EBS volume to all the EC2 instances
- D.Create file systems on Amazon Elastic Block Store (Amazon EBS) volumes that are attached to each EC2 instance Synchronize the EBS volumes across the different EC2 instances

Answer : B

QUESTION 1022

A company used an Amazon RDS for MySQL DB instance during application testing. Before terminating the DB instance at the end of the test cycle a solutions architect created two backups. The solutions architect created the first backup by using the mysqldump utility to create a database dump. The solutions architect created the second backup by enabling the final DB snapshot option on RDS termination. The company is now planning for a new test cycle and wants to create a new DB instance from the most recent backup. The company has chosen a MySQL-compatible edition of Amazon Aurora to host the DB instance. Which solutions will create the new DB instance? (Select TWO)

- A.Import the RDS snapshot directly into Aurora
- B.Upload the RDS snapshot to Amazon S3 then import the RDS snapshot into Aurora
- C.Upload the database dump to Amazon S3 Then import the database dump into Aurora.
- D.Use AWS Database Migration Service (AWS DMS) to import the RDS snapshot into Aurora.
- E.Upload the database dump to Amazon S3 Then use AWS Database Migration Service (AWS DMS) to import the database dump into Aurora

Answer : AC

QUESTION 1023

A medical company is designing a new application that gathers symptoms from patients. The company has decided to use Amazon Simple Queue Service (Amazon SQS) and Amazon Simple Notification Service (Amazon SNS) in the architecture. A solutions architect is reviewing the infrastructure design. Data must be encrypted while at rest and in transit. Only authorized personnel of the company can access the data. Which combination of steps should the solutions architect take to meet these requirements? (Select TWO) .

- A.Turn on server-side encryption on the SQS components Update the default key policy to restrict key usage to a set of authorized principals
- B.Turn on server-side encryption on the SNS components by using a custom CMK Apply a key policy to restrict key usage to a set of authorized principals
- C.Turn on encryption on the SNS components Update the default key policy to restrict key usage to a set of authorized principals Set a condition in the topic policy to allow only encrypted connections over TLS.
- D.Turn on server-side encryption on the SQS components by using a custom CMK. Apply a key policy to restrict key usage to a set of authorized principals Set a condition in the queue policy to allow only encrypted connections over TLS.
- E.Turn on server-side encryption on the SQS components by using a custom CMK. Apply an IAM policy to restrict key usage to a set of authorized principals Set a condition in the queue policy to allow only encrypted connections over TLS.

Answer : CD

QUESTION 1024

A company is building a disaster recovery (DR) solution. The company wants to rotate its primary systems between AWS Regions on a regular basis. The company's application is geographically distributed and includes a serverless web tier. The application's database tier runs on Amazon Aurora. A solutions architect needs to build an architecture for the database layer to implement managed, planned failover. Which combination of actions will meet these requirements with the LEAST downtime? (Select TWO)

- A.Create an Aurora DB cluster Configure Aurora Replicas
- B.Fail over to one of the secondary DB clusters from another Region

- C.Create an Aurora DB cluster snapshot Restore from the snapshot
- D.Configure an Aurora global database Set up a secondary DB cluster
- E.Promote one of the read replicas as a writer from the Amazon RDS console

Answer : CD

QUESTION 1025

A company runs an application on several Amazon EC2 instances that store persistent data on an Amazon Elastic File System (Amazon EFS) file system. The company needs to replicate the data to another AWS Region by using an AWS managed service solution. Which solution will meet these requirements MOST cost-effectively ?

- A.Use the EFS-to-EFS backup solution to replicate the data to an EFS file system in another Region
- B.Run a nightly script to copy data from the EFS file system to an Amazon S3 bucket
Enable S3 Cross-Region Replication on the S3bucket
- C.Create a VPC in another Region Establish a cross-Region VPC peer Run a nightly rsync to copy data from the original Region to the new Region.
- D.Use AWS Backup to create a backup plan with a rule that takes a daily backup and replicates it to another Region Assign the EFS file system resource to the backup plan

Answer : D

QUESTION 1026

A doctor's office is moving all of its patient data to the AWS Cloud. The office needs to retain all the data indefinitely, but the data is rarely accessed after a year. The data must be immediately available during the first year. However, to minimize cost, the office is willing to wait a day for data that is more than 1 year old to become available. Which combination of actions should a solutions architect take to meet these requirements MOST cost-effectively ? (Select TWO)

- A.Create an Amazon S3 Lifecycle transition rule to move the data to S3 Glacier after a year
- B.Create an Amazon S3 Lifecycle transition rule to move the data to S3 Glacier Deep Archive after a year
- C.Create an Amazon S3 bucket for the data. Store data in the S3 bucket by using the S3 Glacier storage class
- D.Create an Amazon S3 bucket for the data. Store data in the bucket by using the S3

Standard storage class.

E.Create an Amazon S3 bucket for the data.Store data in the bucket by using the S3 Intelligent-Tiering storage class.

Answer : AD

QUESTION 1027

A social media company wants to allow its users to upload images in an application that is hosted in the AWS Cloud.The company needs a solution that automatically resizes the images so that the images can be displayed on multiple devicetypes.The application experiences unpredictable traffic patterns throughout the day.The company is seeking a highly available solution that maximizes scalability.What should a solutions architect do to meet these requirements ?

- A.Create a static website hosted in Amazon S3 that invokes AWS Lambda functions to resize the images and store the images in an Amazon S3 bucket
- B.Create a static website hosted in Amazon CloudFront that invokes AWS Step Functions to resize the images and store the images in an Amazon RDS database
- C.Create a dynamic website hosted on a web server that runs on an Amazon EC2 instance Configure a process that runs on the EC2 instance to resize the images and store the images in an Amazon S3 bucket.
- D.Create a dynamic website hosted on an automatically scaling Amazon Elastic Container Service (Amazon ECS) cluster that creates a resize job in Amazon Simple Queue Service (Amazon SQS) Set up an image-resizing program that runs on an Amazon EC2 instance to process the resize jobs.

Answer : A

QUESTION 1028

A company runs batch processes on Amazon EC2 instances that are needed only during business hours.These processes must preserve the data at all times but the speed of processing is not important.The company needs to run these processes in the MOST cost-effective manner.Which solution will meet these requirements ?

- A.Use EC2 Reserved Instances with the All Upfront payment option
- B.Use EC2 Reserved instances with the Partial Upfront payment option
- C.Use Spot Fleet requests with the allocation strategy set to lowestPnace
- D.Use persistent Spot Instance requests with behaviour that stops interrupted instances

Answer : A

QUESTION 1029

A company is creating a three-tier web application consisting of a web server an application server and a database server. The application will track GPS coordinates of packages as they are being delivered. The application will update the database every 0-5 seconds. The tracking will need to be read as fast as possible for users to check the status of their packages. Only a few packages might be tracked on some days whereas millions of packages might be tracked on other days. Tracking will need to be searchable by tracking ID customer ID and order ID Orders older than 1 month no longer need to be tracked. What should a solutions architect recommend to accomplish this with minimal total cost of ownership ?

- A.Use Amazon DynamoDB Enable Auto Scaling on the DynamoDB table Schedule an automatic deletion script for items older than 1month
- B.Use Amazon DynamoDB with global secondary indexes Enable Auto Scaling on the DynamoDB table and the global secondary indexes Enable TTL on the DynamoDB Table
- C.Use an Amazon RDS On-Demand Instance with Provisioned IOPS (PIOPS) Enable Amazon Cloud Watch alarms to send notifications when PIOPS are exceeded Increase and decrease PIOPS as needed
- D.Use an Amazon RDS Reserved Instance with Provisioned IOPS (PIOPS) Enable Amazon Cloud Watch alarms to send notifications when PIOPS are exceeded Increase and decrease PIOPS as needed

Answer : B

QUESTION 1030

A company has a remote factory that has unreliable connectivity. The factory needs to gather and process machine data and sensor data so that it can sense products on its conveyor belts and initiate a robotic movement to direct the products to the right location. Predictable low-latency compute processing is essential for the on-premises control systems. Which solution should the factory use to process the data ?

- A.Amazon CloudFront lambda@edge functions
- B.An Amazon EC2 instance that has enhanced networking enabled
- C.An Amazon EC2 instance that uses an AWS Global Accelerator endpoint

D.An Amazon Elastic Block Store (Amazon EBS) volume on an AWS Snowball Edge cluster

Answer : A

淘宝旺旺 服务真诚128
weixin 淘宝旺旺 服务真诚128